

Computer Organization and Architecture

Module 7

Prof. Indranil Sengupta

Dr. Sarani Bhattacharya

Department of Computer Science and Engineering

IIT Kharagpur

Arithmetic Pipeline

Basic Concept

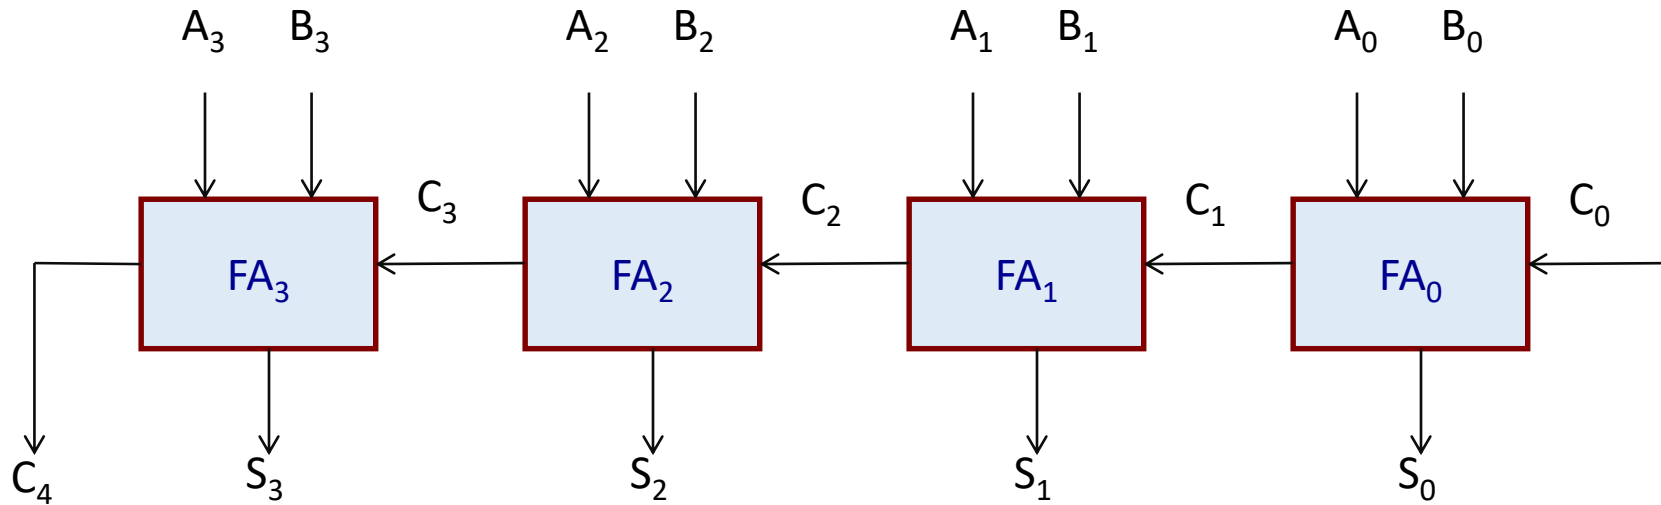
- Various arithmetic operations require a set of simpler computations to be carried out in sequence.
 - Can be split into stages with buffers in between, and run in a pipeline.
- Useful when several similar calculations are required to be carried out in sequence.
 - Example: Vector operations.

```
for (i=0;i<64;i++)  
    A[i] = B[i] * C[i];
```

Fixed Point Addition Pipeline

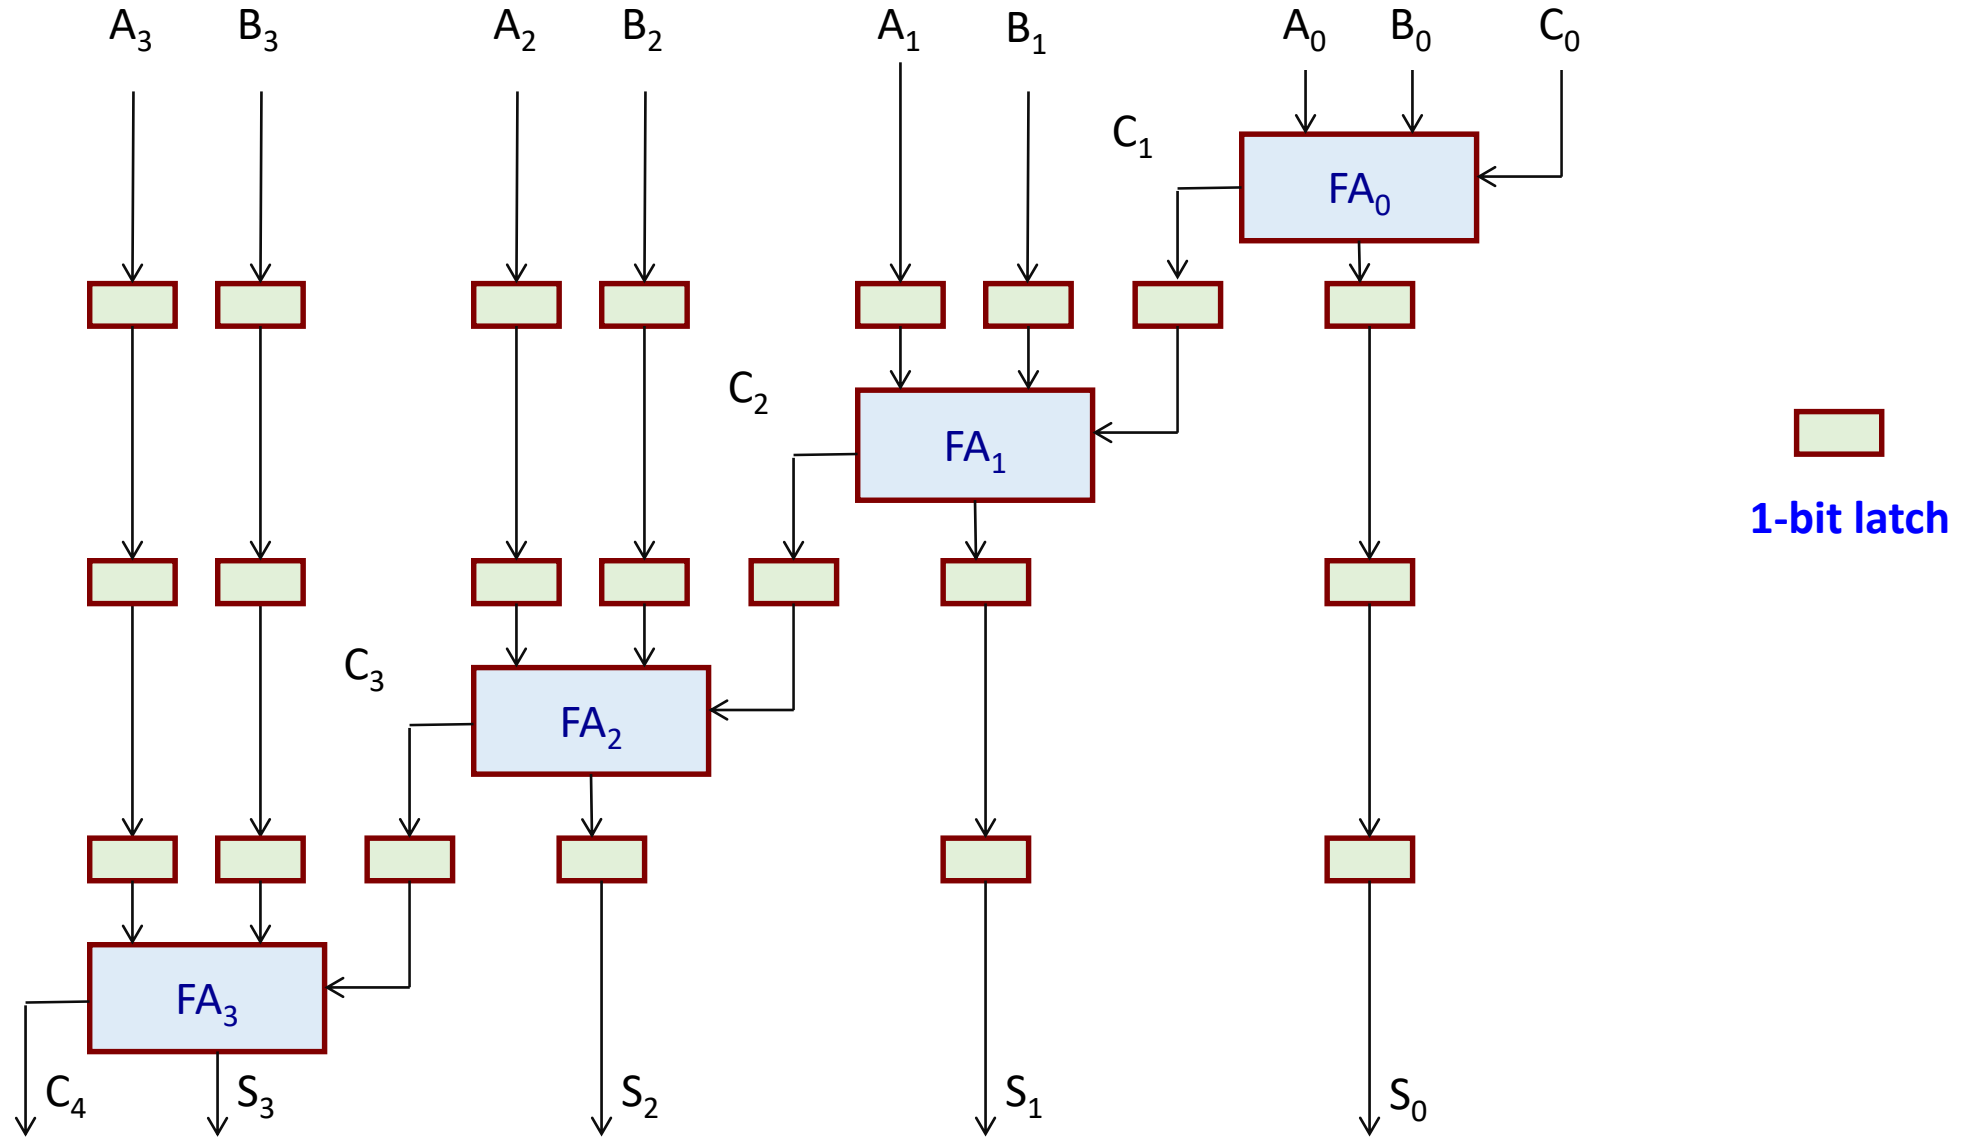
- We have seen how a ripple-carry adder works.
 - Rippling of the carries gives it a bad worst-case performance.
- We explore whether pipelining can improve the performance.
- *Assumption*: delay of a latch is comparable to the delay of a full adder.

A 4-bit Ripple Carry Adder



Worst-case delay $\approx 4 \times$ (carry generation time in FA)

4-bit pipelined ripple-carry adder



- Delay of a full adder = t_{FA}
- Delay of a 1-bit latch = t_L
- Clock period $T \geq (t_{FA} + t_L)$
- After the pipeline is full, one result (sum) is generated every time T .
 - Convenient for vector addition kind of applications.

```
for (i=0; i<10000; i++)  
    a[i] = b[i] + c[i];
```

Floating-Point Addition

- Floating-point addition requires the following steps:
 - a) Compare exponents and align mantissas.
 - b) Add mantissas.
 - c) Normalize result.
 - d) Adjust exponent.
- Subtraction is similar.

Example:

$$A = 0.9504 \times 10^3$$

$$B = 0.8200 \times 10^2$$

Align mantissa: 0.0820

Add mantissa: $0.9504 + 0.0820 = 1.0324$

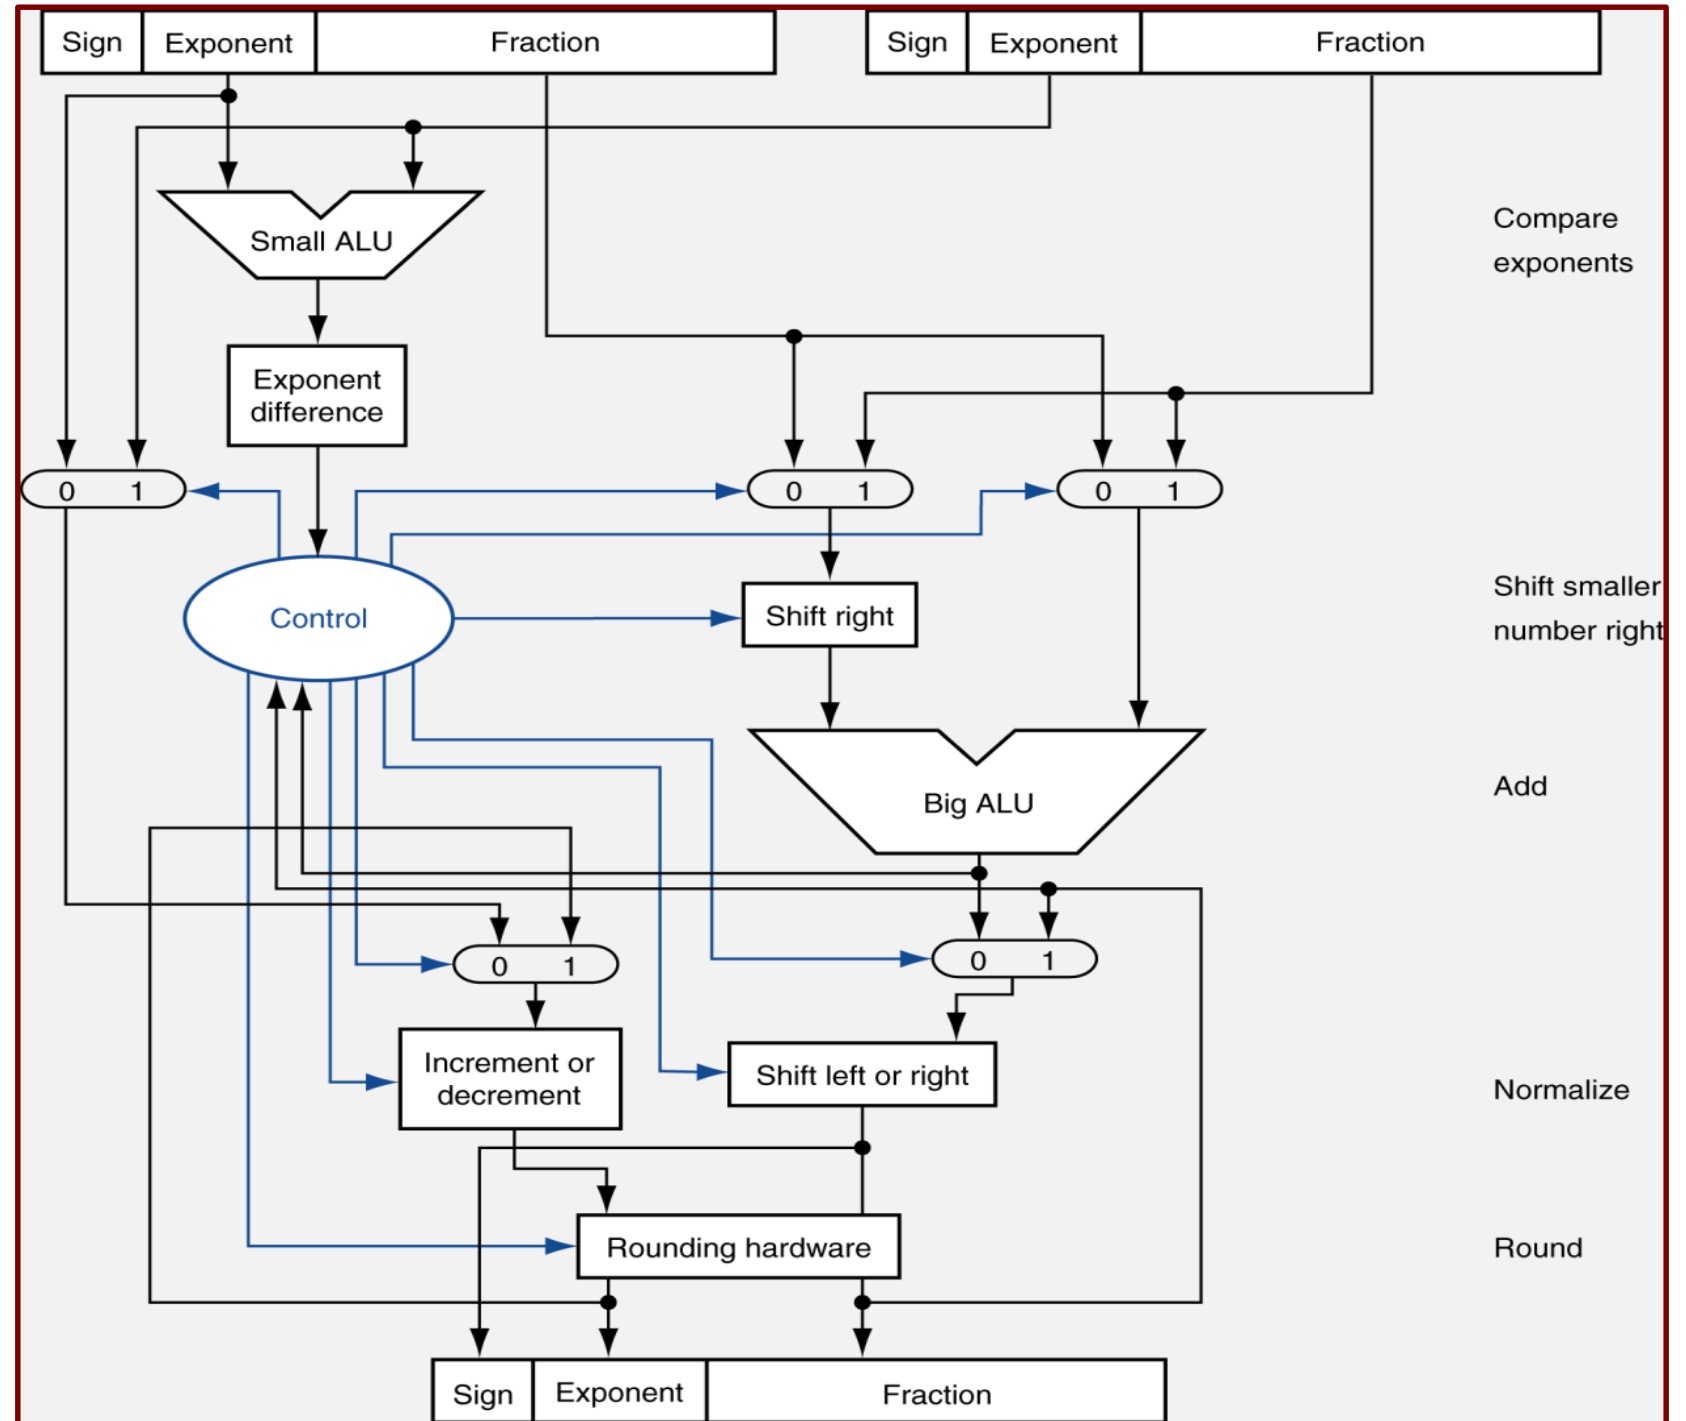
Normalize: 0.10324

Adjust exponent: $3 + 1 = 4$

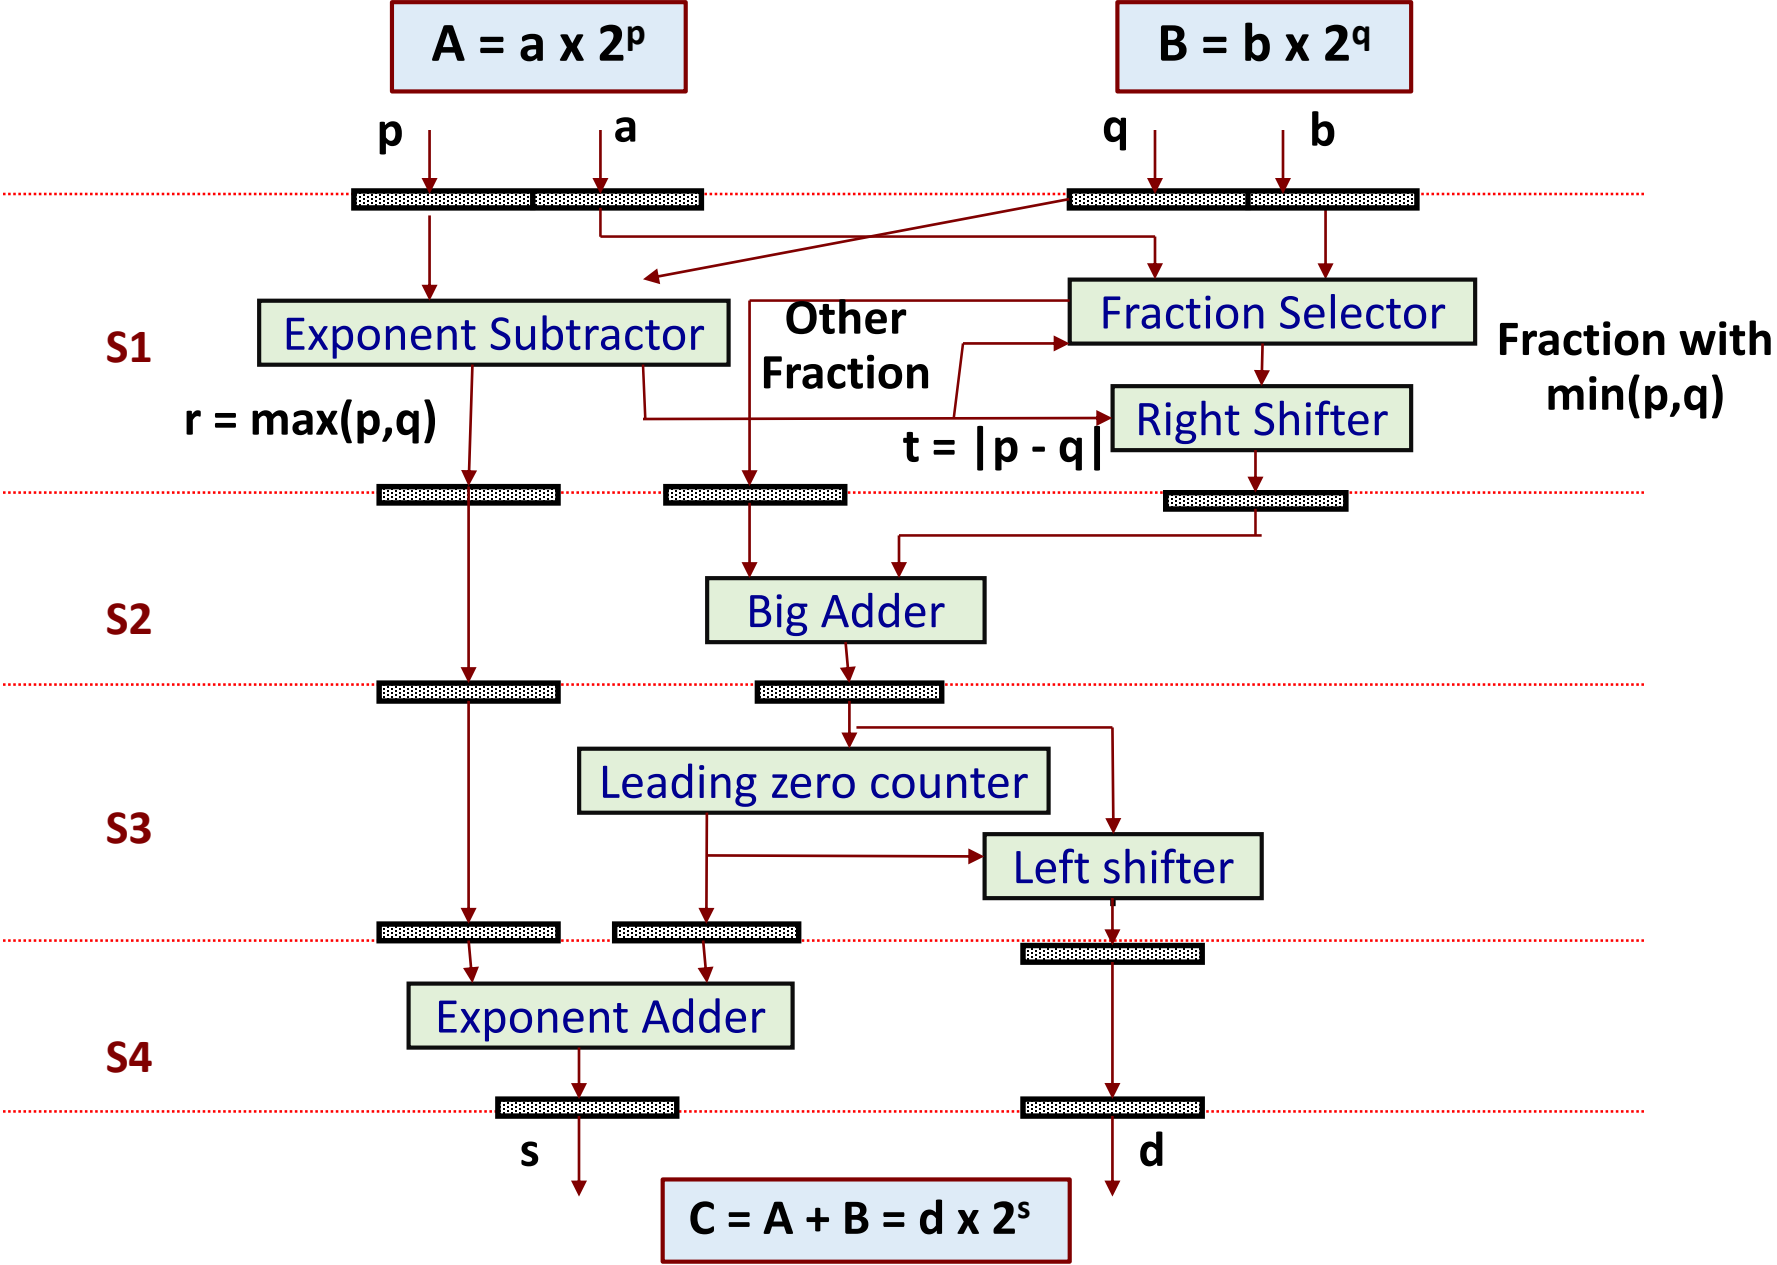
Sum = 0.10324×10^4

Floating-Point Addition Hardware

- The last step of rounding is required in IEEE-754 format.



4-Stage Floating Point Adder



Floating-Point Multiplication

- Floating-point multiplication requires the following steps:
 - a) Add exponents.
 - b) Multiply mantissas.
 - c) Normalize result.
- Division is similar.

A last step of rounding is required in IEEE-754 format.

Example:

$$A = 0.9504 \times 10^3$$

$$B = 0.8200 \times 10^2$$

Add exponents: $3 + 2 = 5$

Multiply mantissa: $0.9504 \times 0.8200 = 0.7793$

Normalize: 0.7793 (no change)


Product = 0.7793×10^5

Pipelining the MIPS32 Data Path

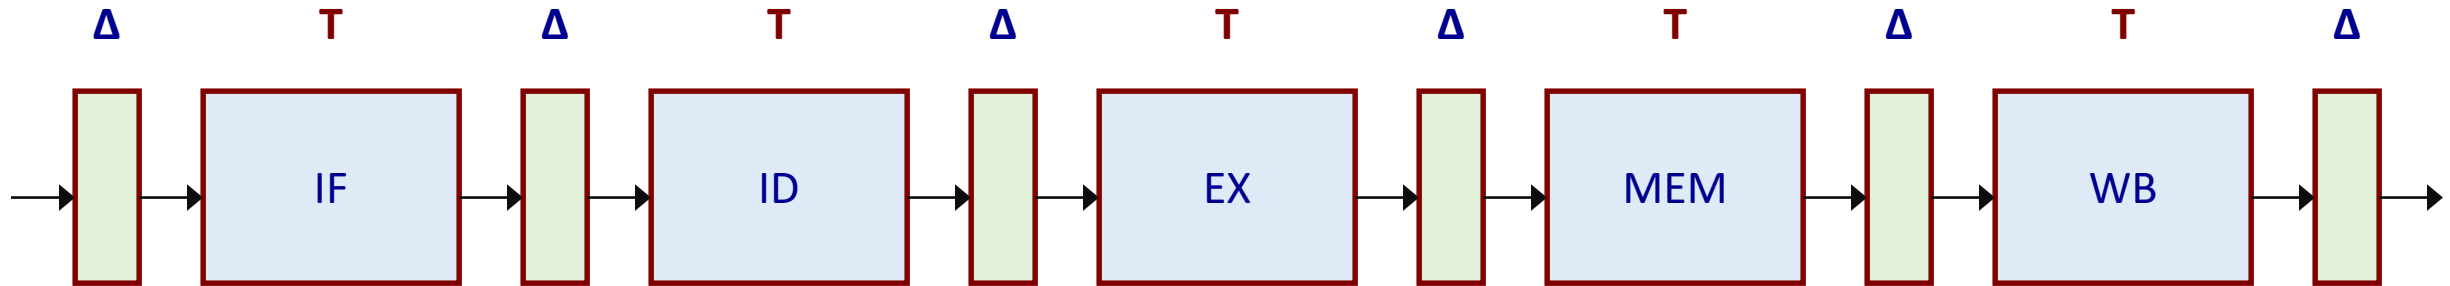
Introduction

- Basic requirements for pipelining the MIPS32 data path:
 - We should be able to start a new instruction every clock cycle.
 - Each of the five steps mentioned before (IF, ID, EX, MEM and WB) becomes a pipeline stage.
 - Each stage must finish its execution within one clock cycle.
- Since execution of several instructions are overlapped, we must ensure that there is no conflict during the execution.
 - Simplicity of the MIPS32 instruction set makes this evaluation quite easy.
 - We shall discuss these issues in some detail.

5T



Time of execute n instructions = $5Tn$



Time of execute n instructions = $(4 + n).(T + \Delta) \approx (4 + n).T$, if $T \gg \Delta$

Ideal Speedup = $5Tn / (4 + n)T \approx 5$, for large n .

In practice, due to various conflicts, speedup is much less.

Clock Cycles

Instruction	1	2	3	4	5	6	7	8
<i>i</i>	IF	ID	EX	MEM	WB			
<i>i + 1</i>		IF	ID	EX	MEM	WB		
<i>i + 2</i>			IF	ID	EX	MEM	WB	
<i>i + 3</i>				IF	ID	EX	MEM	WB

↓
Instr-i
finishes

↓
Instr-(i+1)
finishes

↓
Instr-(i+2)
finishes

↓
Instr-(i+3)
finishes

Clock Cycles

Instruction	1	2	3	4	5	6	7	8
<i>i</i>	IF	ID	EX	MEM	WB			
<i>i + 1</i>		IF	ID	EX	MEM	WB		
<i>i + 2</i>			IF	ID	EX	MEM	WB	
<i>i + 3</i>				IF	ID	EX	MEM	WB

Some examples of conflict:


- IF & MEM: In clock cycle 4, both instructions *i* and *i+3* access memory.
 - *Solution: use separate instructions and data cache.*
- ID & WB: In clock cycle 5, both instructions *i* and *i+3* access register bank.
 - *Solution: allow both read and write access to registers in the same clock cycle.*

Advantages of Pipelining

- In the non-pipelined version, the execution time of an instruction is equal to the combined delay of the five stages (say, $5T$).
- In the pipelined version, once the pipeline is full, one instruction gets executed after every T time.
 - Assuming all state delays are equal (equal to T), and neglecting latch delay.
- However, due to various conflicts between instructions (called *hazards*), we cannot achieve the ideal performance.
 - Several techniques have been proposed to improve the performance.
 - To be discussed.


Some Observations

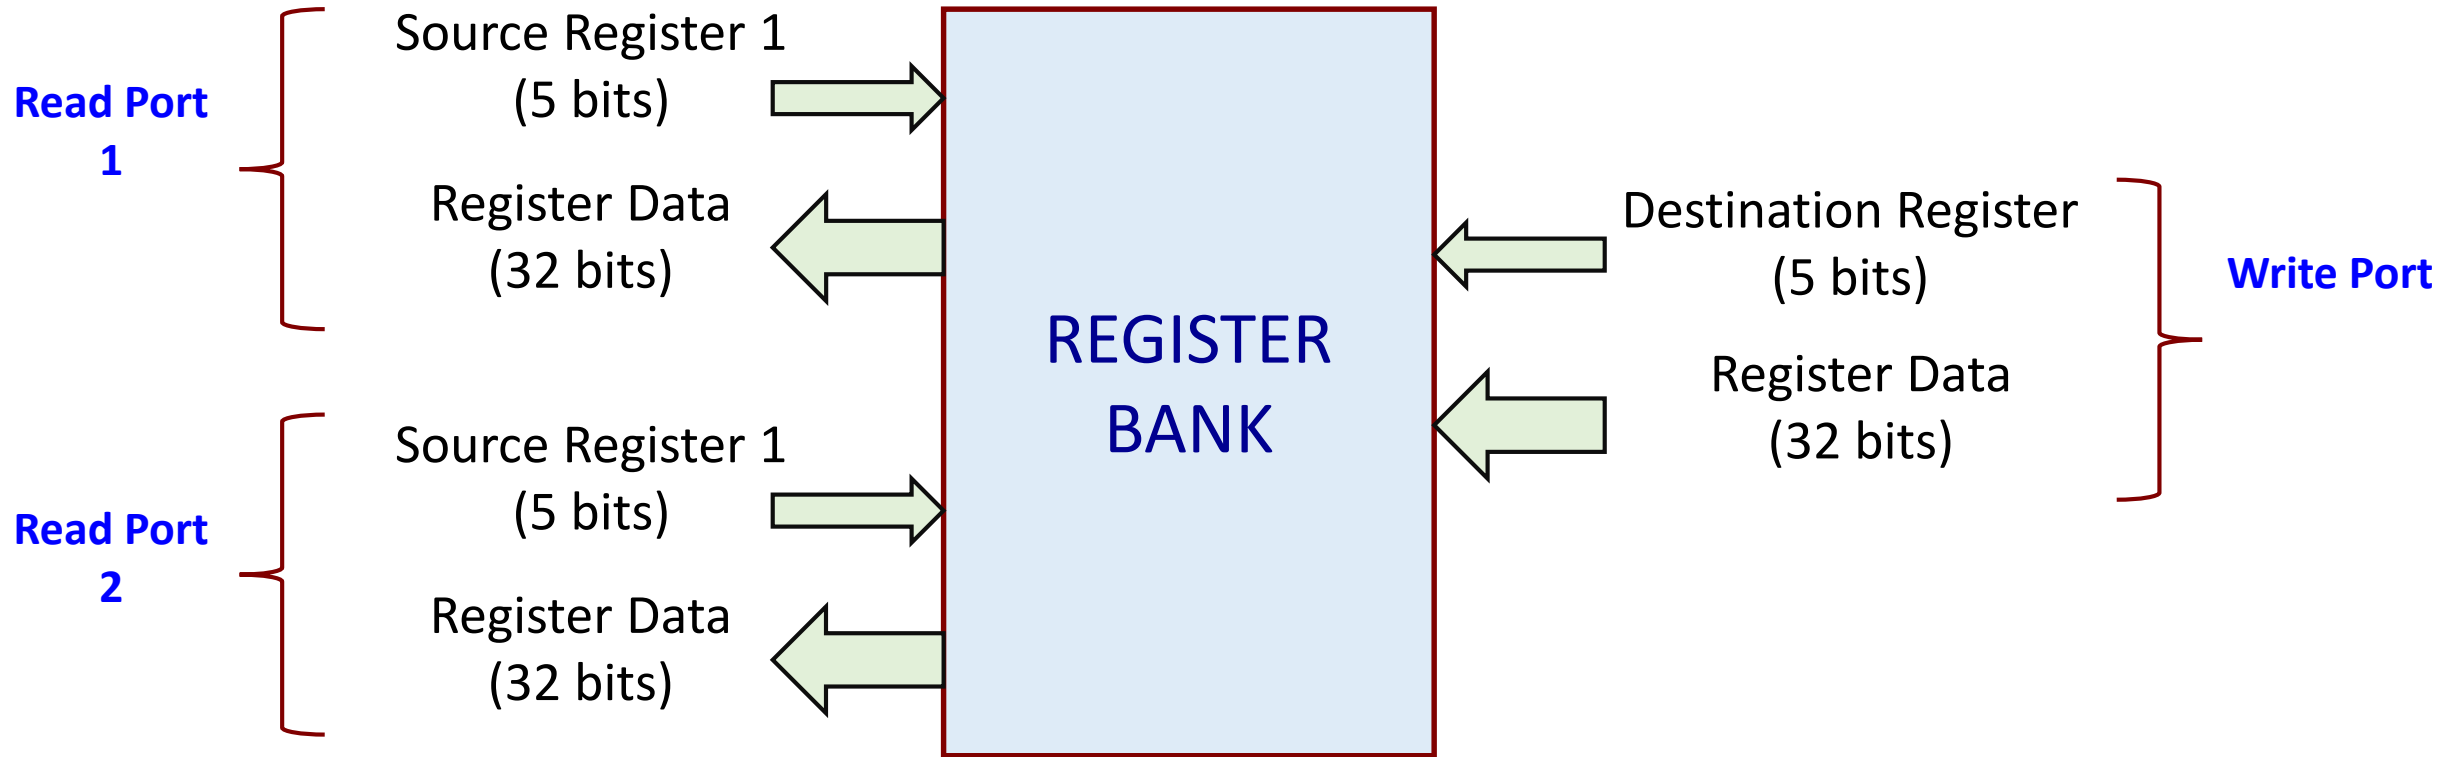
- a) To support overlapped execution, peak memory bandwidth must be increased *5 times* over that required for the non-pipelined version.
- An instruction fetch occurs every clock cycle.
 - Also there can be two memory accesses per clock cycle (one for instruction and one for data).
 - Separate instruction and data caches are typically used to support this.



b) The register bank is accessed both in the stages ID and WB.

- ID requires 2 register reads, and WB requires 1 register write.
- We thus have the requirement of *2 reads and 1 write* in every clock cycle.
- Two register reads can be supported by having two register read ports.
- Simultaneous reads and write may result in clashes (e.g., same register used).
- Solution adopted in MIPS32 pipeline is to perform the write during the *first half* of the clock cycle, and the reads during the *second half* of the clock cycle.





- c) Since a new instruction is fetched every clock cycle, it is required to increment the *PC* on each clock.
- PC updating has to be done *during IF stage itself*, as otherwise the next instruction cannot be fetched.
 - In the non-pipelined version discussed earlier, this was done during the MEM stage.

Basic Performance Issues in a Pipeline



- Register stages are inserted between pipeline stages, which increases the execution time of an individual instruction.
 - Because of overlapped execution of instructions, throughput increases.
- The clock period T has to be chosen suitably:
 - Slowest stage in the pipeline.
 - Clock skew and jitter.
 - Register setup time: minimum time the register input must be held stable before the active clock edge arrives.

Example 1

- Consider the 5-stage MIPS32 pipeline, with the following features:
 - Pipeline clock rate of 1GHz (i.e. 1 ns clock cycle time).
 - For a non-pipelined implementation, ALU operations and branches take 4 cycles, while memory operations take 5 cycles.
 - Relative frequencies of ALU operations, branches and memory operations are 50%, 15%, and 35% respectively.
 - In the pipelined implementation, due to clock skew and setup time, the clock cycle time increases by 0.25 ns.
 - Calculate the estimated speedup of the pipelined implementation.

- **Solution:**

a) For non-pipelined processor:

- Average instruction execution time = Clock cycle time x Average CPI
$$= 1 \text{ ns} \times (0.50 \times 4 + 0.15 \times 4 + 0.35 \times 5) = 4.35 \text{ ns}$$

b) For pipelined processor:

- Clock cycle time = $1 + 0.25 = 1.25 \text{ ns}$
- In the steady state, one instruction will get executed every clock cycle.
- Speedup = $4.35 / 1.25 = 3.48$

Micro-operations for Non-pipelined MIPS32

IF

```
IR ← Mem [PC];  
NPC ← PC + 4;
```

ID

```
A ← Reg [rs];  
B ← Reg [rt];  
Imm ← (IR15)16 ## IR15..0  
Imm1 ← IR25..0 ## 00
```

EX

```
ALUOut ← A + Imm;
```

(a) Memory

```
ALUOut ← A func B;
```

(b) R-R ALU

```
ALUOut ← A func Imm;
```

(c) R-IMM ALU

```
ALUOut ← NPC + (Imm << 2);  
cond ← (A op 0);
```

(d) Branch

MEM

```
PC ← NPC;  
LMD ← Mem [ALUOut];
```

(a) Load

```
PC ← NPC;  
Mem [ALUOut] ← B;
```

(b) Store

```
if (cond) PC ← ALUOut;  
else PC ← NPC;
```

(c) Branch

```
PC ← NPC;
```

(d) Others

WB

Reg [rd] \leftarrow ALUOut;

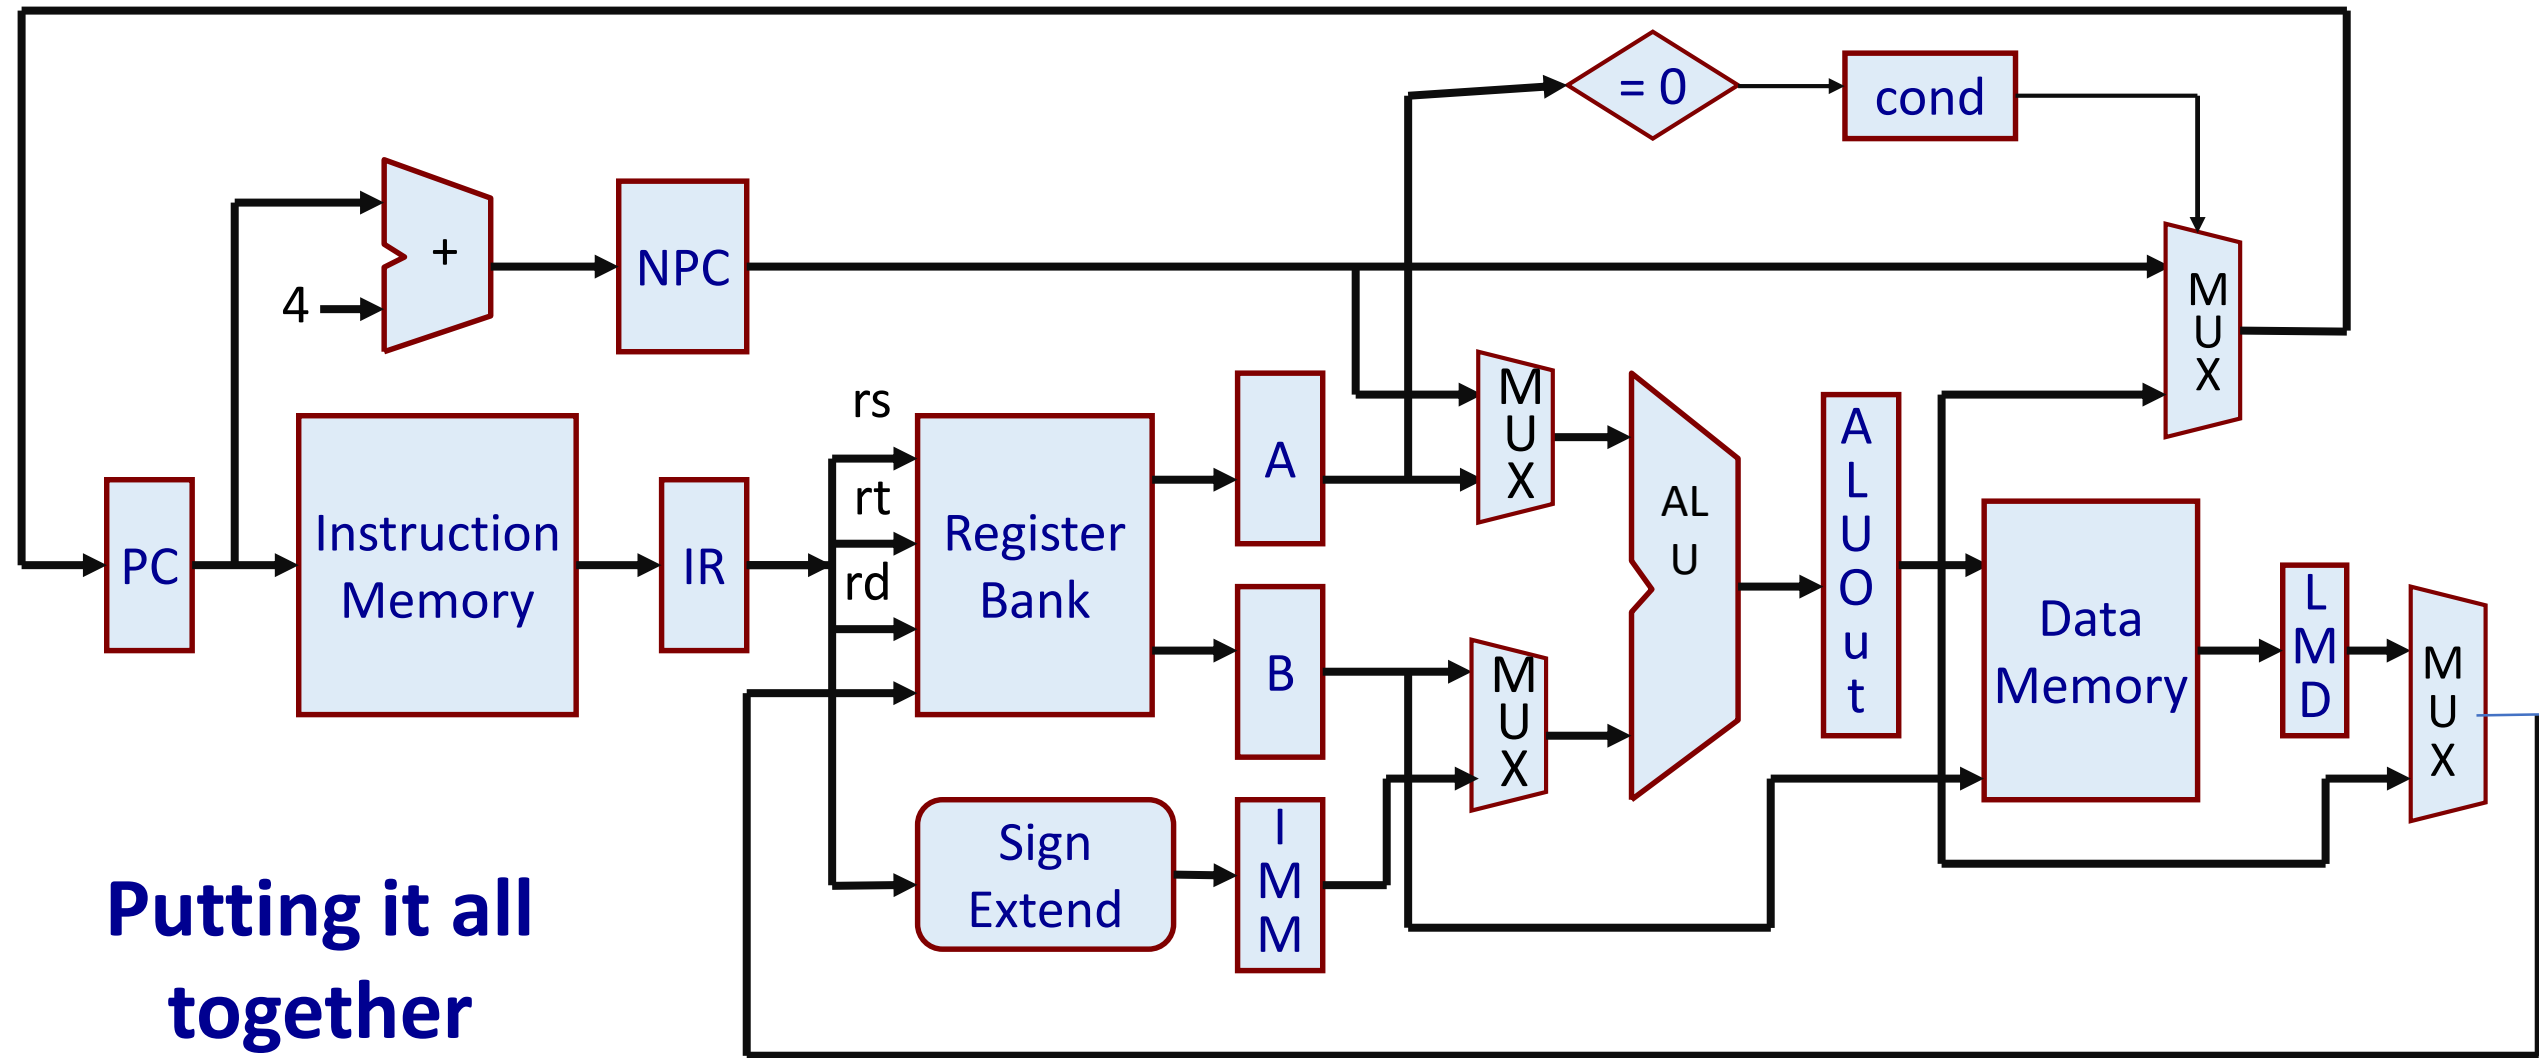
(a) R-R ALU

Reg [rt] \leftarrow ALUOut;

(b) R-IMM ALU

Reg [rt] \leftarrow LMD;

(c) Load

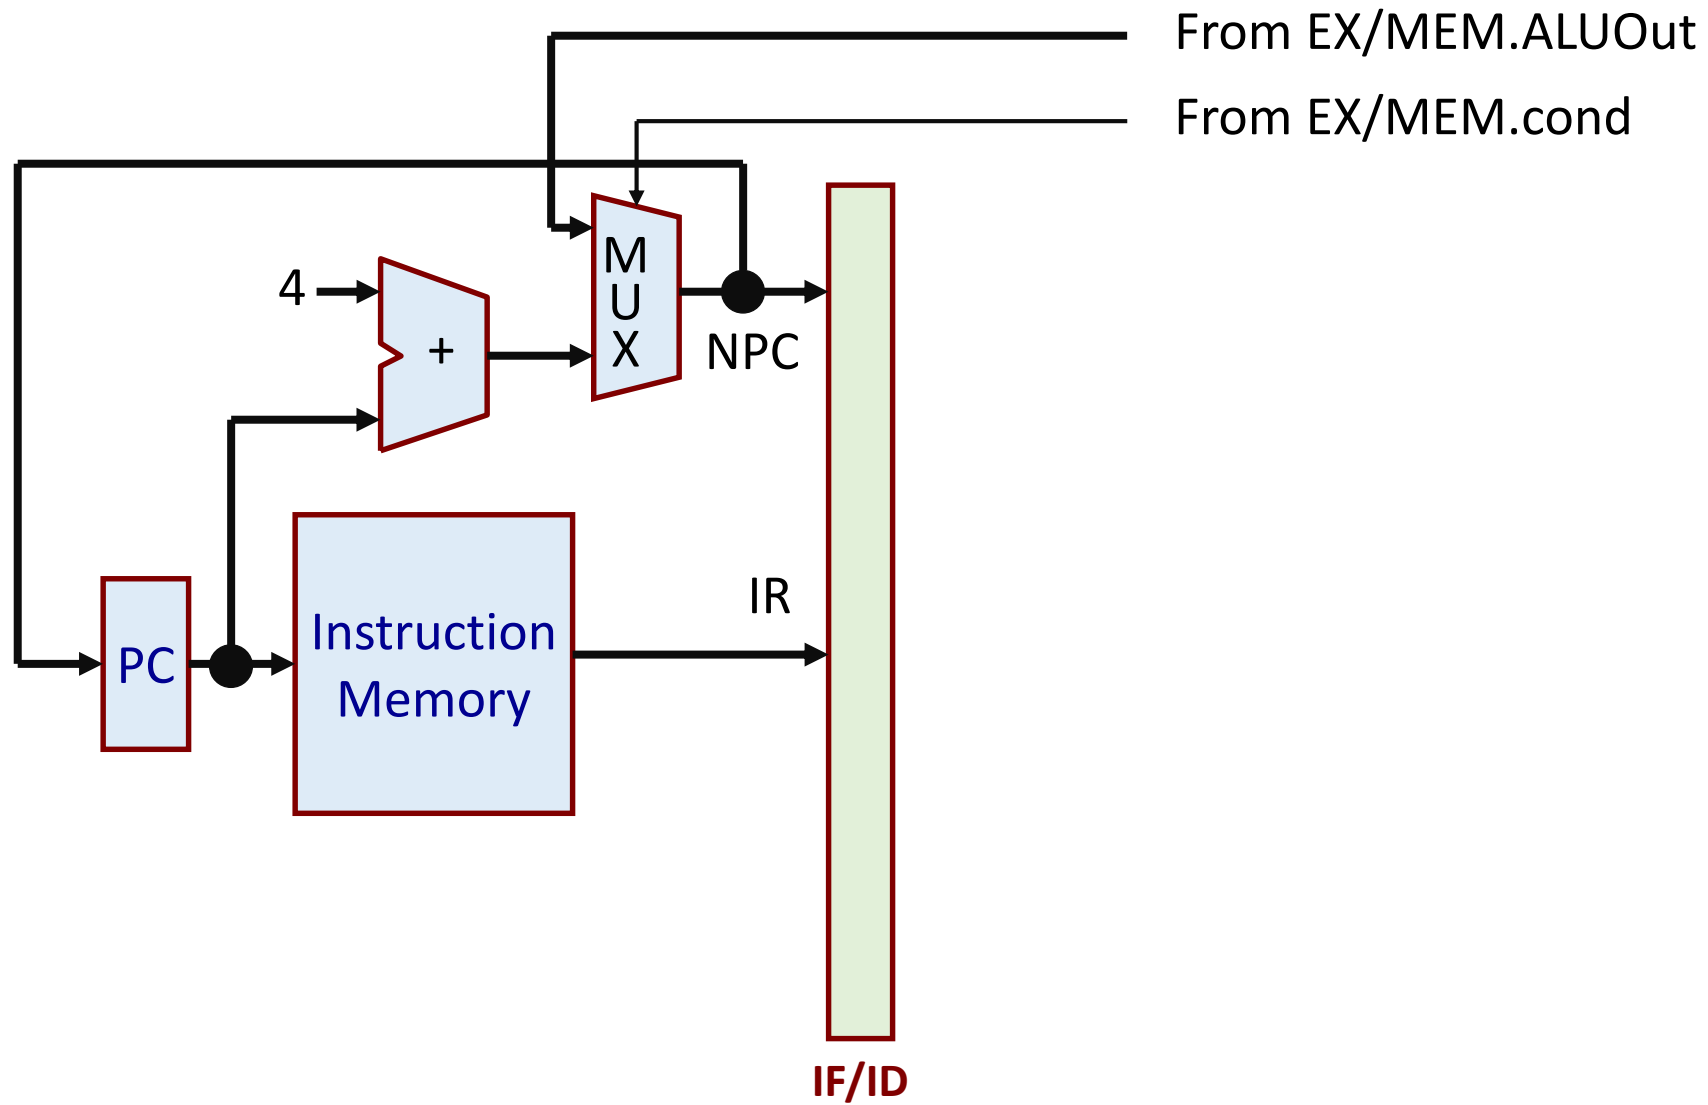


Micro-operations for Pipelined MIPS32

- Convention used:
 - Many of the temporary registers required in the data path are included as part of the inter-stage latches.
 - IF/ID: denotes the latch stage between the IF and ID stages.
 - ID/EX: denotes the latch stage between the ID and EX stages.
 - EX/MEM: denotes the latch stage between the EX and MEM stages.
 - MEM/WB: denotes the latch stage between the MEM and WB stages.
- **Example:**
 - ID/EX.A means a register *A* that is implemented as part of the *ID/EX latch* stage.

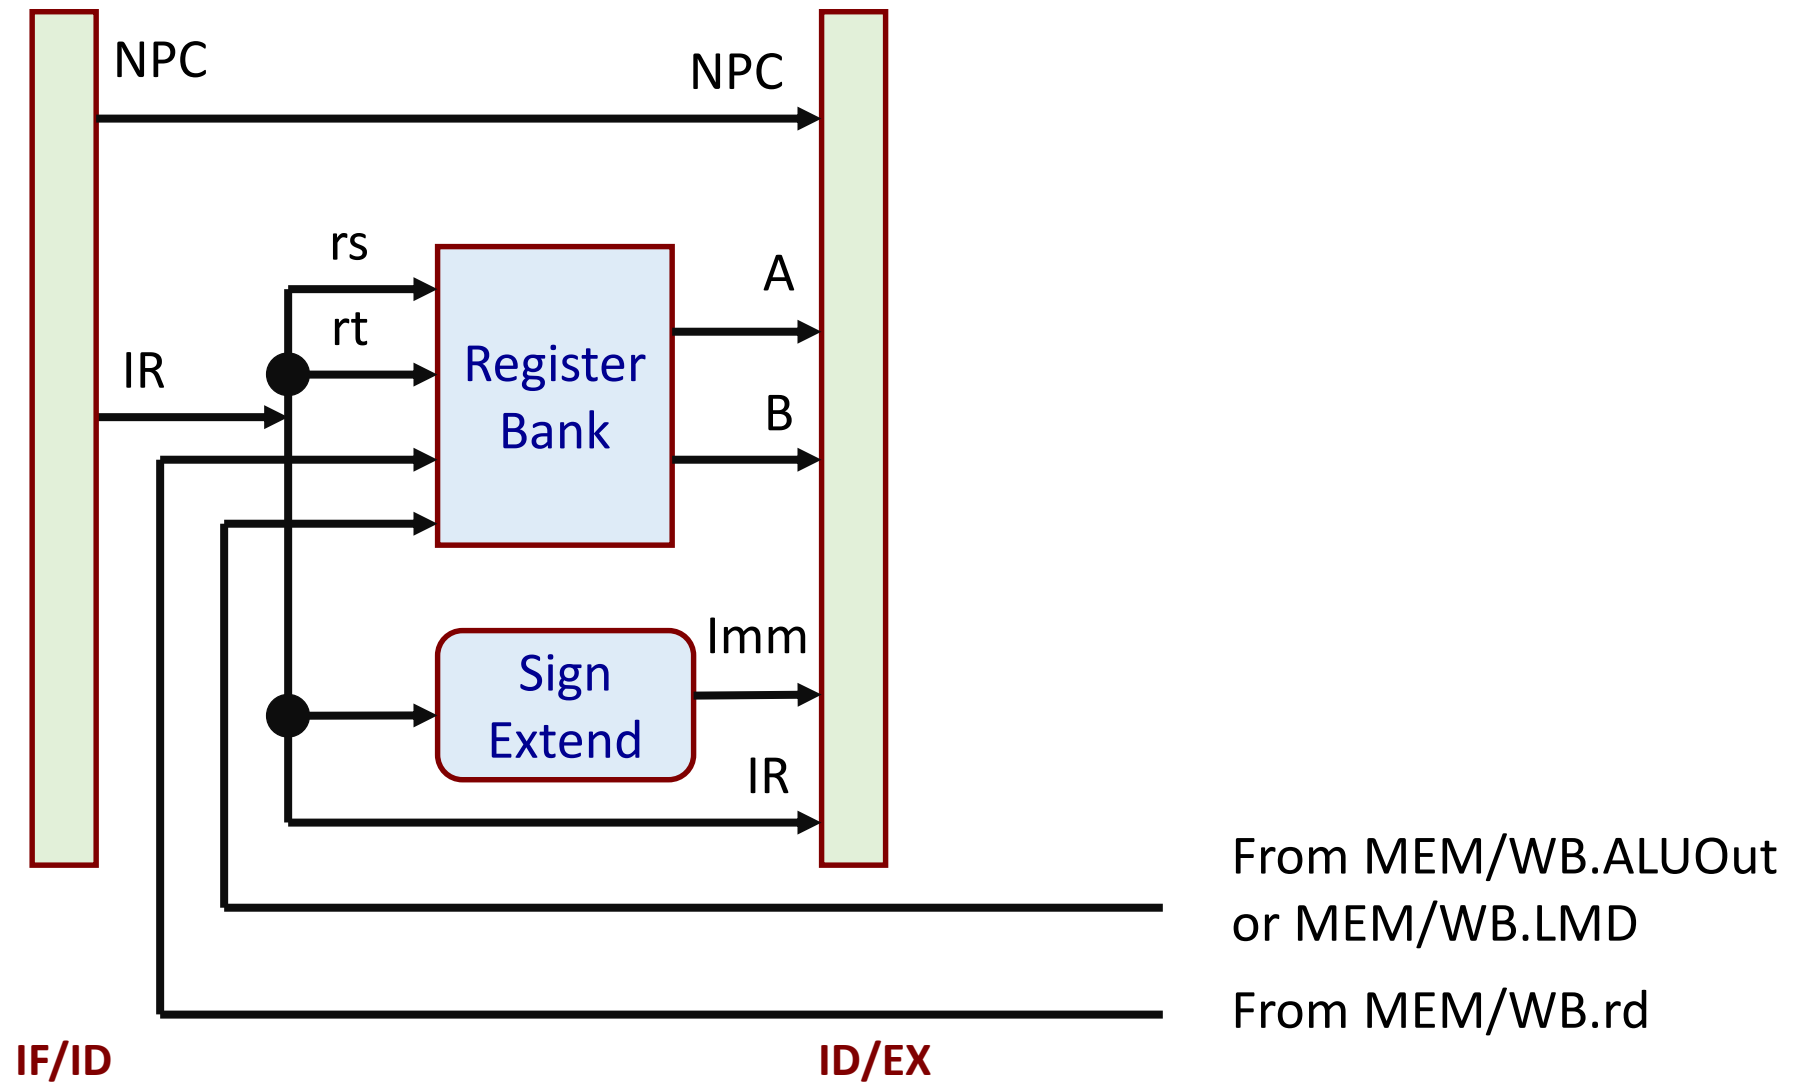
(a) Micro-operations for Pipeline Stage IF

```
IF/ID.IR      ← Mem [PC];  
IF/ID.NPC,PC ← ( if ((EX/MEM.opcode == branch) & EX/MEM.cond)  
                { EX/MEM.ALUOut}  
                else {PC + 4} );
```



(b) Micro-operations for Pipeline Stage ID

ID/EX.A \leftarrow Reg [IF/ID.IR [rs]];
ID/EX.B \leftarrow Reg [IF/ID.IR [rt]];
ID/EX.NPC \leftarrow IF/ID.NPC;
ID/EX.IR \leftarrow IF/ID.IR;
ID/EX.Imm \leftarrow sign-extend (IF/ID.IR_{15..0});



(c) Micro-operations for Pipeline Stage EX

```
EX/MEM.IR ← ID/EX.IR;
EX/MEM.ALUOut ← ID/EX.A func ID/EX.B;
```

R-R ALU

```
EX/MEM.IR ← ID/EX.IR;  
EX/MEM.ALUOut ← ID/EX.A func ID/EX.Imm;
```


R-M ALU

```
EX/MEM.IR ← ID/EX.IR;  
EX/MEM.ALUOut ← ID/EX.A + ID/EX.B;  
EX/MEM.B ← ID/EX.B;
```

LOAD / STORE

```
EX/MEM.ALUOut ← ID/EX.NPC +  
                (ID.EX.Imm << 2);  
EX/MEM.cond ← (ID/EX.A == 0);
```

BRANCH



(d) Micro-operations for Pipeline Stage MEM

MEM/WB.IR \leftarrow EX/MEM.IR;
MEM/WB.ALUOut \leftarrow EX/MEM.ALUOut;

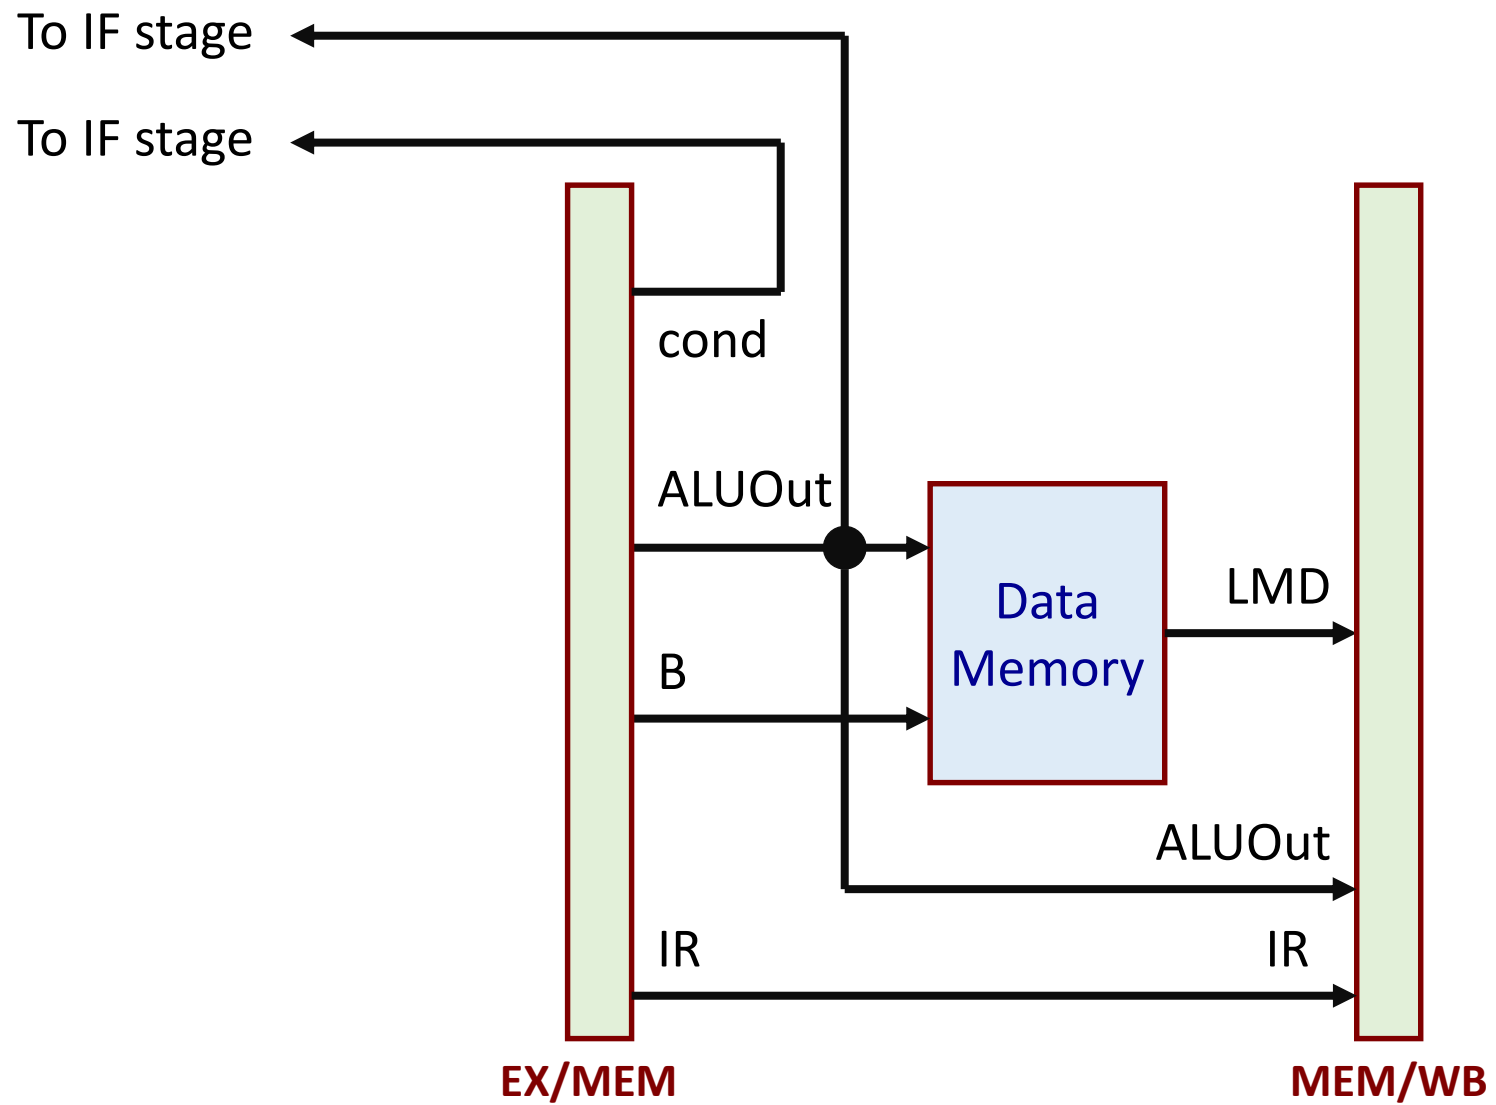
ALU

MEM/WB.IR \leftarrow EX/MEM.IR;
MEM/WB.LMD \leftarrow Mem [EX/MEM.ALUOut];

LOAD

MEM/WB.IR \leftarrow EX/MEM.IR;
Mem [EX/MEM.ALUOut] \leftarrow EX/MEM.B;

STORE

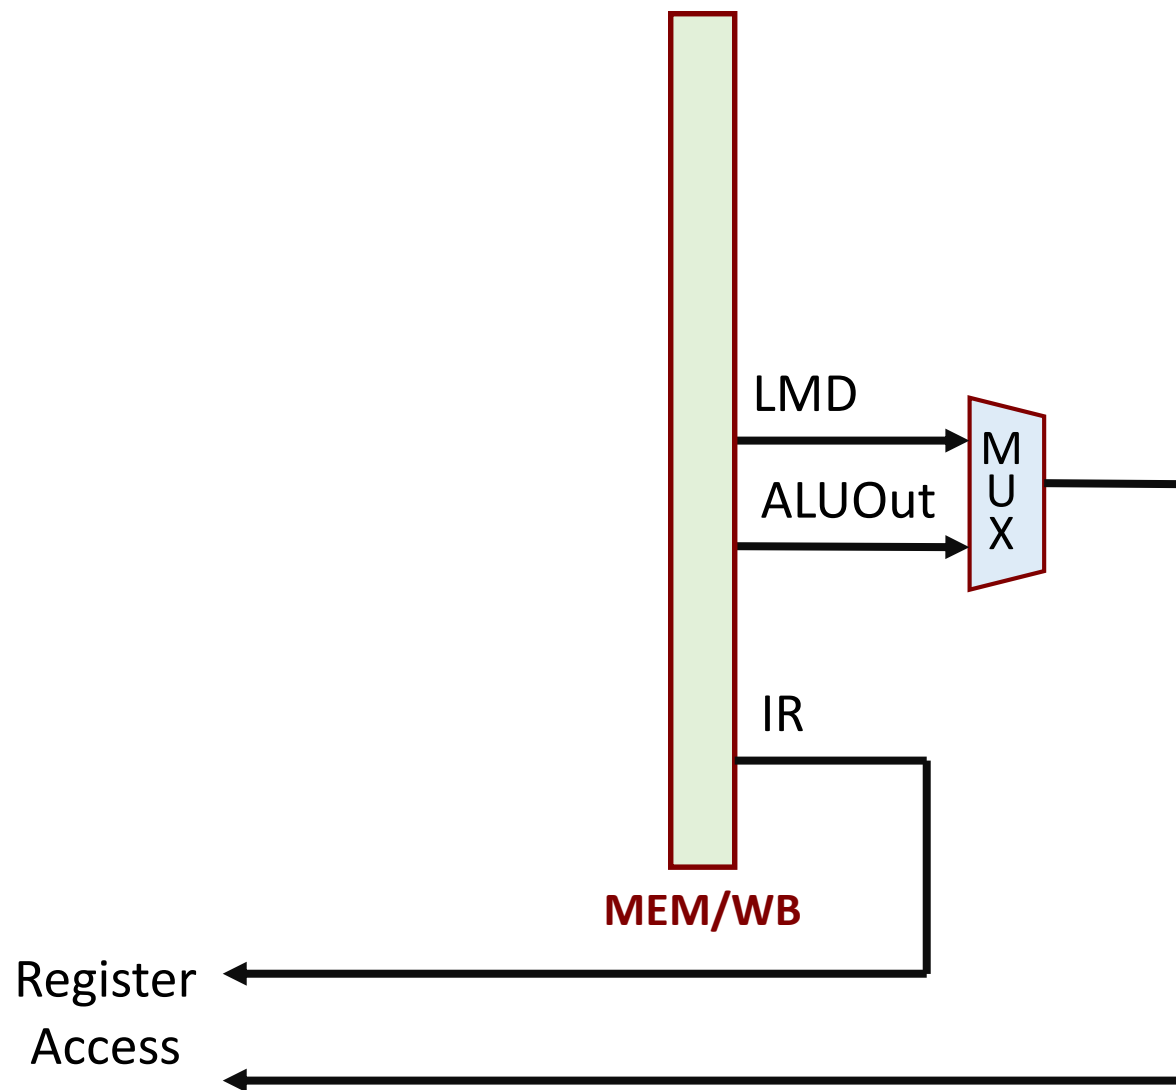


(e) Micro-operations for Pipeline Stage WB

Reg [MEM/WB.IR [rd]] \leftarrow MEM/WB. ALUOut; **R-R ALU**

Reg [MEM/WB.IR [rt]] \leftarrow MEM/WB. ALUOut; **R-M ALU**

Reg [MEM/WB.IR [rt]] \leftarrow MEM/WB. LMD; **LOAD**



PUTTING IT ALL TOGETHER :: MIPS32 PIPELINE

