

Computer Organization and Architecture

Module 4- Part 2

Prof. Indranil Sengupta

Dr. Sarani Bhattacharya

Department of Computer Science and Engineering

IIT Kharagpur

Memory Hierarchy Design

Introduction

- Programmers want unlimited amount of memory with very low latency.
- Fast memory technology is more expensive per bit than slower memory.
 - SRAM is more expensive than DRAM, DRAM is more expensive than disk.
- Possible solution?
 - Organize the memory system in several levels, called *memory hierarchy*.
 - Exploit temporal and spatial locality on computer programs.
 - Try to keep the commonly accessed segments of program / data in the faster memories.
 - Results in faster access times on the average.

Quick Review of Memory Technology

- **Static RAM:**

- Very fast but expensive memory technology (requires 6 transistors / bit).
- Packing density is limited.

- **Dynamic RAM:**

- Significantly slower than SRAM, but much less expensive (1 transistor / bit).
- Requires periodic refreshing.

- **Flash memory:**

- Non-volatile memory technology that uses floating-gate MOS transistors.
- Slower than DRAM, but higher packing density, and lower cost per bit.

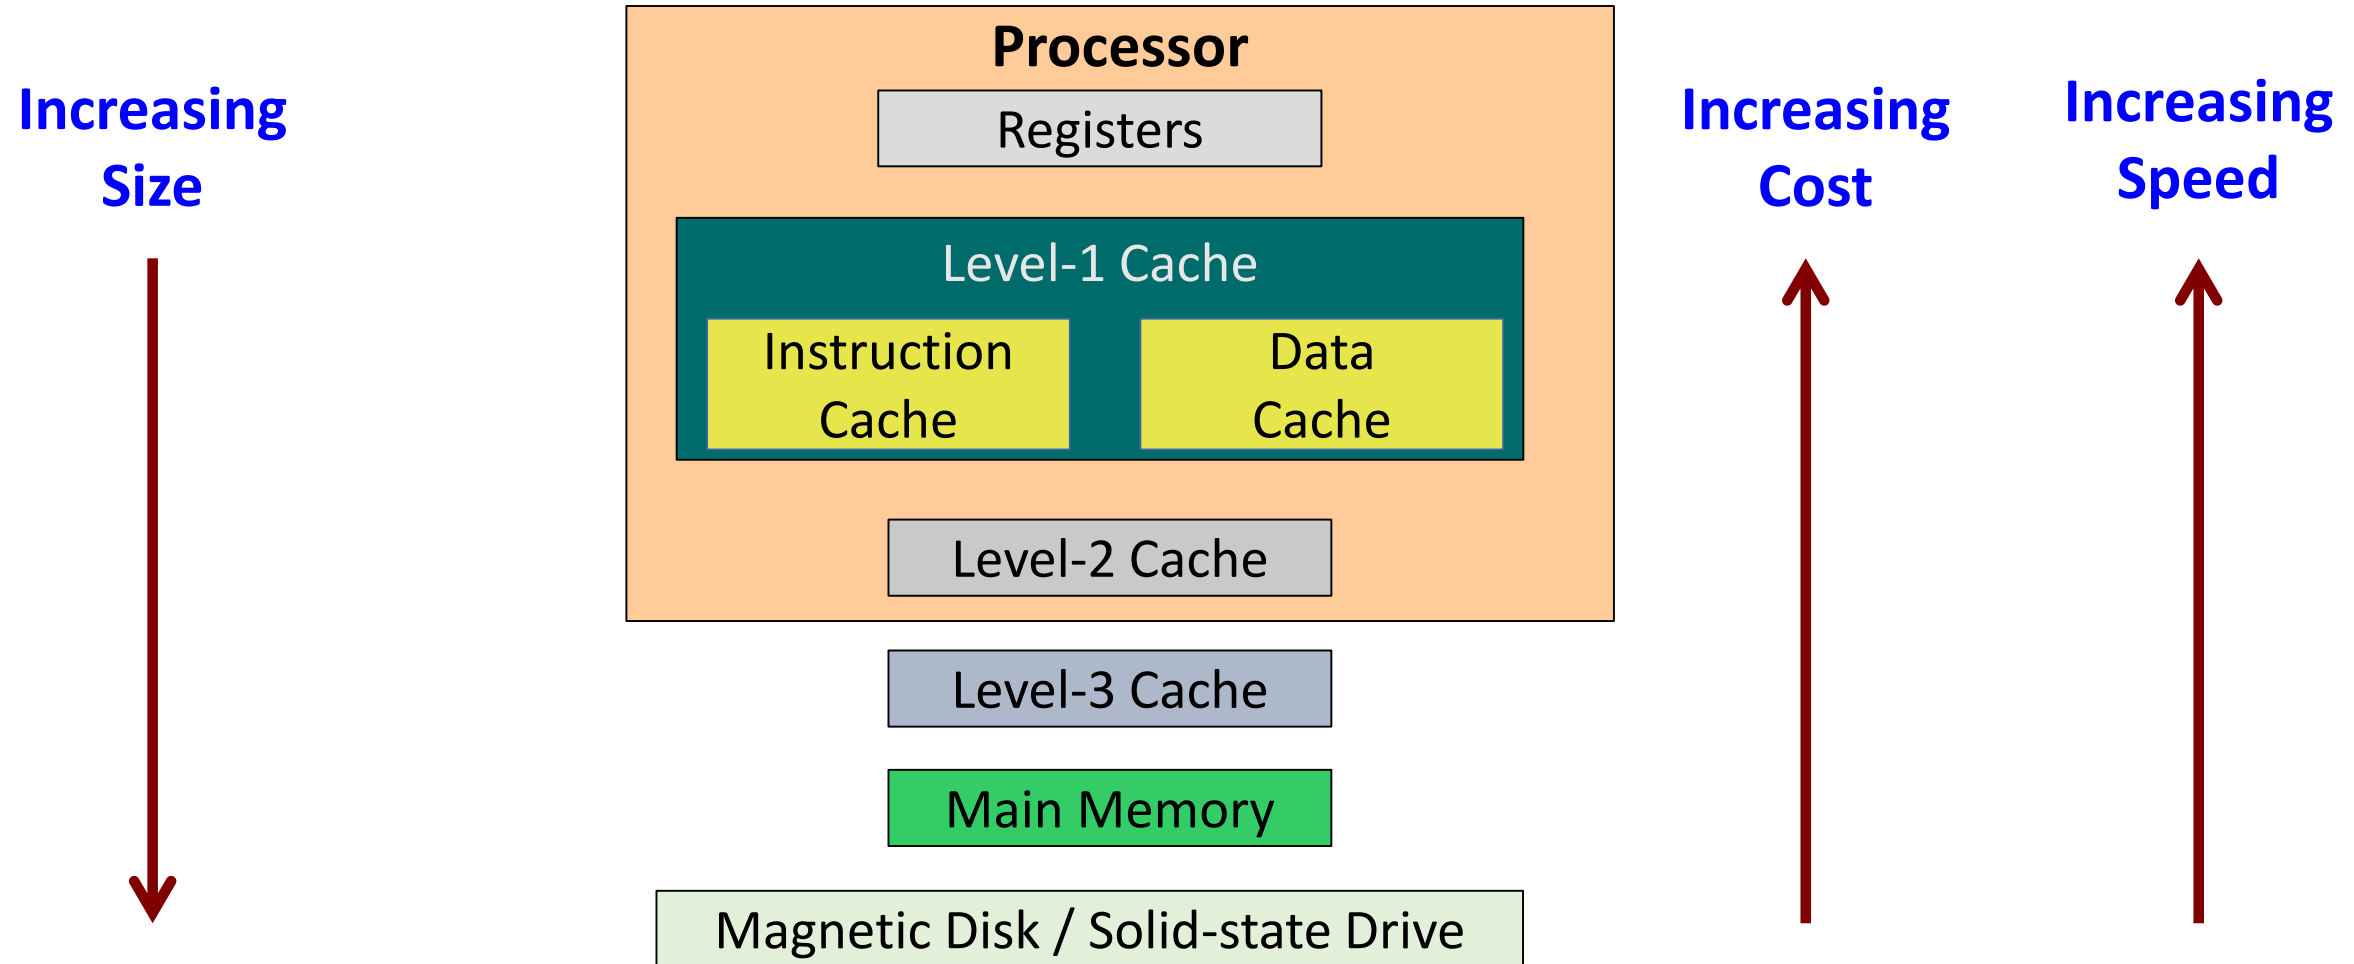
- **Magnetic disk:**

- Provides large amount of storage, with very low cost per bit.
- Much slower than DRAM, and also flash memory.
- Requires mechanical moving parts, and uses magnetic recording technology.

Memory Hierarchy

- The memory system is organized in several levels, using progressively faster technologies as we move towards the processor.
 - The entire addressable memory space is available in the largest (but slowest) memory (typically, magnetic disk or flash storage).
 - We incrementally add smaller (but faster) memories, each containing a *subset* of the data stored in the memory below it.
 - We proceed in steps towards the processor.

- Typical hierarchy (starting with closest to the processor):
 1. Processor registers
 2. Level-1 cache (typically divided into separate instruction and data cache)
 3. Level-2 cache
 4. Level-3 cache
 5. Main memory
 6. Secondary memory (magnetic disk / solid-state drive)
- As we move away from the processor:
 - a) Size increases
 - b) Cost decreases
 - c) Speed decreases




A Comparison

Level	Typical Access Time	Typical Capacity	Other Features
Register	300-500 ps	500-1000 B	On-chip
Level-1 cache	1-2 ns	16-64 KB	On-chip
Level-2 cache	5-20 ns	256 KB – 2 MB	On-chip
Level-3 cache	20-50 ns	1-32 MB	On or off chip
Main memory	50-100 ns	1-16 GB	
Magnetic disk	5-50 ms	100 GB – 16 TB	

Processor-Memory Performance Gap

- Processor is much faster than main memory.
 - Processor has to spend much of the time waiting while instructions and data are being fetched from main memory.
- Memory speed cannot be increased beyond a certain point.



Impact of Processor / Memory Performance Gap

Year	CPU Clock	Clock Cycle	Memory Access	Minimum CPU Stall Cycles
1986	8 MHz	125 ns	190 ns	$190 / 125 - 1 = 0.5$
1989	33 MHz	30 ns	165 ns	$165 / 30 - 1 = 4.5$
1992	60 MHz	16.6 ns	120 ns	$120 / 16.6 - 1 = 6.2$
1996	200 MHz	5 ns	110 ns	$110 / 5 - 1 = 21.0$
1998	300 MHz	3.33 ns	100 ns	$100 / 3.33 - 1 = 29.0$
2000	1 GHz	1 ns	90 ns	$90 / 1 - 1 = 89.0$
2002	2 GHz	0.5 ns	80 ns	$80 / 0.5 - 1 = 159.0$
2004	3 GHz	0.33 ns	60 ns	$60 / 0.33 - 1 = 179.0$

Ideal memory access
time = 1 CPU cycle

Real memory access
time \gg 1 CPU cycle

- **Memory Latency Reduction Techniques:**

- Faster DRAM cells (depends on VLSI technology)
- Wider memory bus width (fewer memory accesses needed)
- Multiple memory banks
- Integration of memory controller with processor
- New emerging RAM technologies

- **Memory Latency Hiding Techniques**

- Memory hierarchy (using SRAM-based cache memories)
- Pre-fetching instructions and/or data from memory before they are actually needed (used to hide long memory access latency)

Locality of Reference


- Programs tend to reuse data and instructions they have used recently.
 - **Rule of thumb:** 90% of the total execution time of a program is spent in only 10% of the code (also called 90/10 rule).
 - **Reason:** nested loops in a program, few procedures calling each other repeatedly, arrays of data items being accessed sequentially, etc.
- Basic idea to exploit this rule:
 - Based on a program's recent past, we can predict with a reasonable accuracy what instructions and data will be accessed in the near future.

- The 90/10 rule has two dimensions:
 - a) **Temporal Locality** (locality in time)
 - If an item is referenced in memory, it will tend to be referenced again soon.
 - b) **Spatial locality** (locality in space)
 - If an item is referenced in memory, nearby items will tend to be referenced soon.

(a) Temporal Locality

- Recently executed instructions are likely to be executed again very soon.
- Example: computing factorial of a number.

```
fact = 1;  
for k = 1 to N  
    fact = fact * k;
```




```
ADDI    $t1,$zero,1  
ADDI    $t2,$zero,N  
ADDI    $t3,$zero,1  
Loop:   MUL    $t1,$t1,$t3  
        ADDI    $t3,$t3,1  
        SGT     $t4,$t3,$t2  
        BNEZ    $t4,Loop
```

- The four instructions in the loop are executed more frequently than the others.

(b) Spatial Locality

- Instructions residing close to a recently executing instruction are likely to be executed soon.
- Example: accessing elements of an array.

```
sum = 0;  
for k = 1 to N  
    sum = sum + A[k];
```



```
        SUB    $t1,$t1,$t1  
        ADDI   $t2,$zero,N  
        ADDI   $t3,$zero,1  
        ADDI   $t5,$zero,A  
Loop:   LW     $t8,0($t5)  
        ADD    $t1,$t1,$t8  
        ADDI   $t3,$t3,1  
        SGT    $t4,$t3,$t2  
        BNEZ   $t4,Loop
```

- Performance can be improved by copying the array into cache memory.


- An example:
 - Accessing a 2-D array row-wise or column-wise.
 - Assume that the array elements are stored row-wise in memory.
 - One row of data can be brought into cache memory at any given time.

```
for (i=0; i<128; i++)  
    for (j=0; j<128; j++)  
        A[i][j] = A[i][j] + 1;
```

```
for (j=0; j<128; j++)  
    for (i=0; i<128; i++)  
        A[i][j] = A[i][j] + 1;
```

Performance of Memory Hierarchy

- We first consider a 2-level hierarchy consisting of two levels of memory, say, M_1 and M_2 .



a) Cost:

- Let c_i denote the cost per bit of memory M_i , and S_i denote the storage capacity in bits of M_i .
- The average cost per bit of the memory hierarchy is given by:

$$\text{Cost } c = \frac{c_1 S_1 + c_2 S_2}{S_1 + S_2}$$

b) Hit Ratio / Hit Rate:

- The hit ratio H is defined as the probability that a logical address generated by the CPU refers to information stored in M_1 .
- We can determine H experimentally as follows:
 - A set of representative programs is executed or simulated.
 - The number of references to M_1 and M_2 , denoted by N_1 and N_2 respectively, are recorded.

$$H = \frac{N_1}{N_1 + N_2}$$

- The quantity $(1 - H)$ is called the *miss ratio*.

c) Access Time:

- Let t_{A1} and t_{A2} denote the access times of M_1 and M_2 respectively, relative to the CPU.
- The average time required by the CPU to access a word in memory can be expressed as:

$$t_A = H.t_{A1} + (1 - H).t_{\text{MISS}}$$

where t_{MISS} denotes the time required to handle the miss, called *miss penalty*.

- The miss penalty t_{MISS} can be estimated in various ways:
 - a) The simplest approach is to set $T_{\text{MISS}} = t_{A2}$, that is, when there is a miss the data is accessed directly from M_2 .
 - b) A request for a word not in M_1 typically causes a block containing the requested word to be transferred from M_2 to M_1 . After completion of the block transfer, the word can be accessed in M_1 .

If t_B denotes the block transfer time, we can write

$$t_{\text{MISS}} = t_B + t_{A1} \quad [\text{since } t_B \gg t_{A1}, t_{A2} \approx t_B]$$

Thus, $t_A = H.t_{A1} + (1 - H).(t_B + t_{A1})$

- c) If t_{HIT} denotes the time required to check whether there is a hit, we can write

$$t_{\text{MISS}} = t_{\text{HIT}} + t_B + t_{A1}$$

d) Efficiency:

- Let $r = t_{A2} / t_{A1}$ denote the access time ratio of the two levels of memory.
- We define the access efficiency as $e = t_{A1} / t_A$, which is the factor by which t_A differs from its minimum possible value.

$$\text{Efficiency } e = \frac{t_{A1}}{H \cdot t_{A1} + (1 - H) \cdot t_{A2}} = \frac{1}{H + (1 - H) \cdot r}$$

e) Speedup:

- The speedup gained by using the memory hierarchy is defined as $S = t_{A2} / t_A$.
- We can write:

$$S = \frac{t_{A2}}{H \cdot t_{A1} + (1 - H) \cdot t_{A2}} = \frac{1}{(1 - H) + H / r}$$

- The same result follows from *Amdahl's law*.

Some Common Terminologies Used

- **Block:** The smallest unit of information transferred between two levels.
- **Hit Rate:** The fraction of memory accesses found in the upper level.
- **Hit Time:** Time to access the upper level
 - Upper level access time + Time to determine hit/miss
- **Miss:** Data item needs to be retrieved from a block in the lower level.
- **Miss Rate:** The fraction of memory accesses not found in the upper level.
- **Miss Penalty:** Overhead whenever a miss occurs.
 - Time to replace a block in the upper level + Time to transfer the missed block

Example 1

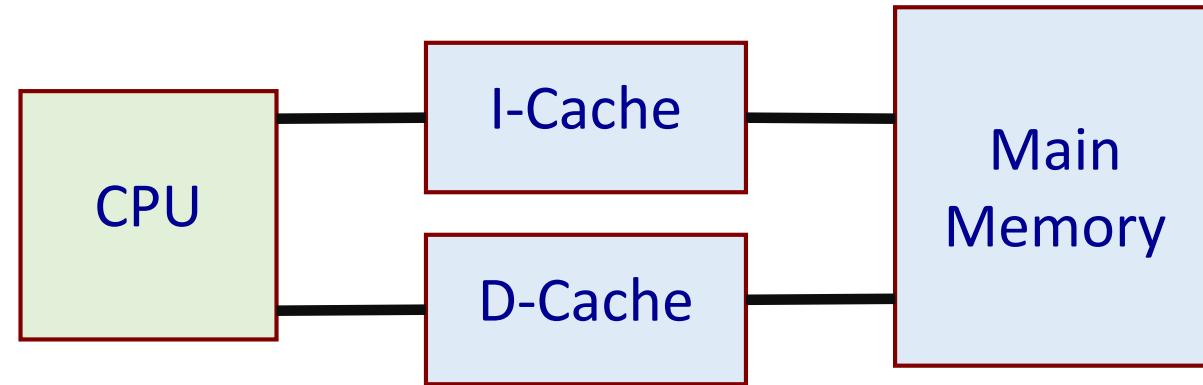
- Consider a 2-level memory hierarchy consisting of a cache memory M_1 and the main memory M_2 . Suppose that the cache is 6 times faster than the main memory, and the cache can be used 90% of the time. How much speedup do we gain by using the cache?

Here, $r = 6$ and $H = 0.90$

Thus, $S = 1 / [H / r + (1 - H)] = 1 / (0.90 / 6 + 0.10) = 1 / 0.25 = 4$

Example 2


- Consider a 2-level memory hierarchy with separate instruction and data caches in level 1, and main memory in level 2.



The following parameters are given:

- The clock cycle time is 2 ns.
- The main mem. access time is 15 clock cycles (for both read and write).
- 1 % of instructions are not found in I-cache.
- 8 % of data references are not found in D-cache.
- 20 % of the total memory accesses are for data.
- Cache access time (including hit detection) is 1 clock cycle.

Determine the overall average access time.



$$t_{\text{MISS}} = 1 + 15 = 16 \text{ cycles}$$

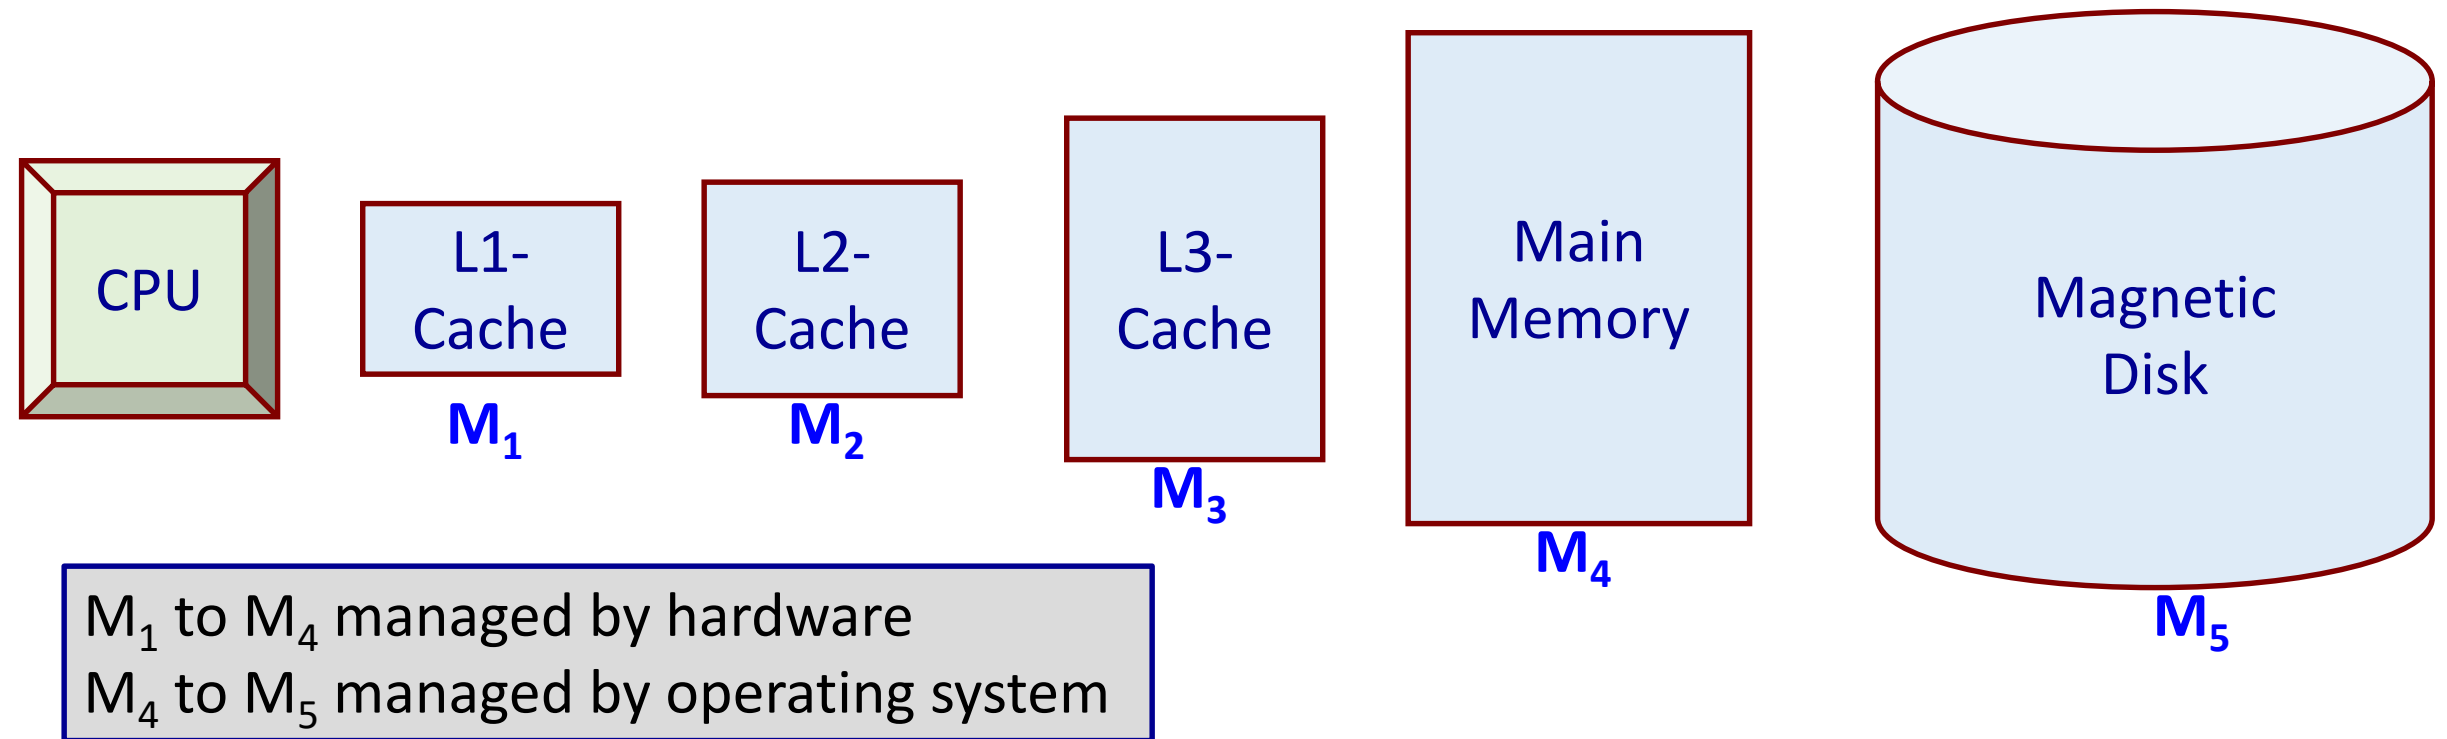
Average number of cycles per access:


$$\begin{aligned} & 0.80 \times (0.99 \times 1 + 0.01 \times 16) + 0.20 \times (0.92 \times 1 + 0.08 \times 16) \\ & = 0.92 + 0.44 = 1.36 \end{aligned}$$

Thus, average access time $t_A = 1.36 \times 2 \text{ ns} = 2.72 \text{ ns}$

Performance Calculation for Multi-Level Hierarchy

- Most of the practical memory systems use more than 2 levels of hierarchy.





t_{L1} : access time of M_1

t_{L2} : access time of M_2

H_{L1} : hit ratio of M_1

H_{L2} : hit ratio of M_2 with respect to the residual accesses that try to access M_2

- Consider a 3-level hierarchy consisting of L1-cache, L2-cache and main memory.
- Whenever there is a miss in L1, we go to L2.
- Average access time can be calculated as:

$$t_A = H_{L1} \cdot t_{L1} + (1 - H_{L1}) \cdot [H_{L2} \cdot t_{L2} + (1 - H_{L2}) \cdot t_{\text{MISS2}}]$$

- Here, t_{MISS2} is the miss penalty when the requested data is found neither in M_1 nor in M_2 .

Implications of a Memory Hierarchy to the CPU

- Processors designed without memory hierarchy are simpler because all memory accesses take the same amount of time.
 - Misses in a memory hierarchy implies *variable memory access times* as seen by the CPU.
- Some mechanism is required to determine whether or not the requested information is present in the top level of the memory hierarchy.
 - Check happens on every memory access and affects hit time.
 - Implemented in hardware to provide acceptable performance.

- Some mechanism is required to transfer blocks between consecutive levels.
 - If the block transfer requires 10's of clock cycles (like in cache / main memory hierarchy), it is controlled by *hardware*.
 - If the block transfer requires 1000's of clock cycles (like in main memory / secondary memory hierarchy), it can be controlled by *software*.
- Four main questions:
 1. *Block Placement*: Where to place a block in the upper level?
 2. *Block Identification*: How is a block found if present in the upper level?
 3. *Block Replacement*: Which block is to be replaced on a miss?
 4. *Write Strategy*: What happens on a write?

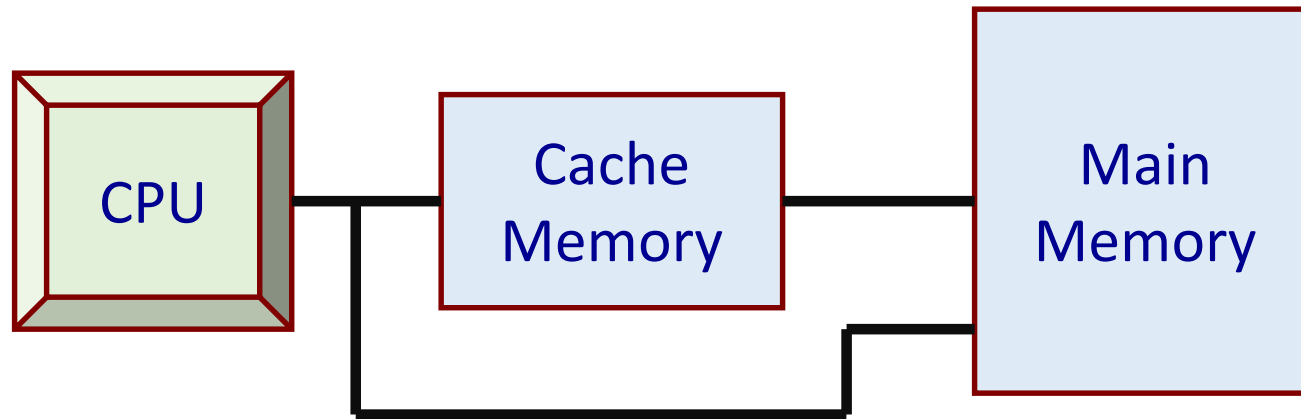
Common Memory Hierarchies

- In a typical computer system, the memory system is managed as two different hierarchies.
 - The *Cache / Main Memory hierarchy*, which consists of 2 to 4 levels and is managed by hardware.
 - Main objective: provide fast average memory access.
 - The *Main Memory / Secondary Memory hierarchy*, which consists of 2 levels and is managed by software (operating system).
 - Main objective: provide large memory space for users (virtual memory).

Cache Memory

Introduction

- Let us consider a single-level cache, and that part of the memory hierarchy consisting of cache memory and main memory.



- Cache memory is logically divided into *blocks* or *lines*, where every block (line) typically contains 8 to 256 bytes.
- When the CPU wants to access a word in memory, a special hardware first checks whether it is present in cache memory.
 - If so (called *cache hit*), the word is directly accessed from the cache memory.
 - If not, the block containing the requested word is brought from main memory to cache.
 - For writes, sometimes the CPU can also directly write to main memory.
- Objective is to keep the commonly used blocks in the cache memory.
 - Will result in significantly improved performance due to the property of *locality of reference*.

Q1. Where can a block be placed in the cache?

- This is determined by some *mapping algorithms*.
 - Specifies which main memory blocks can reside in which cache memory blocks.
 - At any given time, only a small subset of the main memory blocks can be held in main memory.
- Three common block mapping techniques are used:
 - a) Direct Mapping**
 - b) Associative Mapping**
 - c) (N-way) Set Associative Mapping**

An Example Scenario: A 2-level memory hierarchy

- Consider a 2-level cache memory / main memory hierarchy.
 - The cache memory consists of 256 blocks (lines) of 32 words each.
 - Total cache size is 8192 (8K) words.
 - Main memory is addressable by a 24-bit address.
 - Main memory is word addressable.
 - Total size of the main memory is $2^{24} = 16 \text{ M}$ words.
 - Number of 32-word blocks in main memory = $16 \text{ M} / 32 = 512\text{K}$

(a) Direct Mapping

- Each main memory block can be placed in only one block in the cache.

- The mapping function is:

$$\text{Cache Block} = (\text{Main Memory Block}) \% (\text{Number of cache blocks})$$


- For the example,

$$\text{Cache Block} = (\text{Main Memory Block}) \% 256$$


- Some example mappings:

$0 \rightarrow 0, 1 \rightarrow 1, 255 \rightarrow 255, 256 \rightarrow 0, 257 \rightarrow 1, 512 \rightarrow 0, 513 \rightarrow 1, \text{ etc.}$


Direct Mapping



Cache Memory

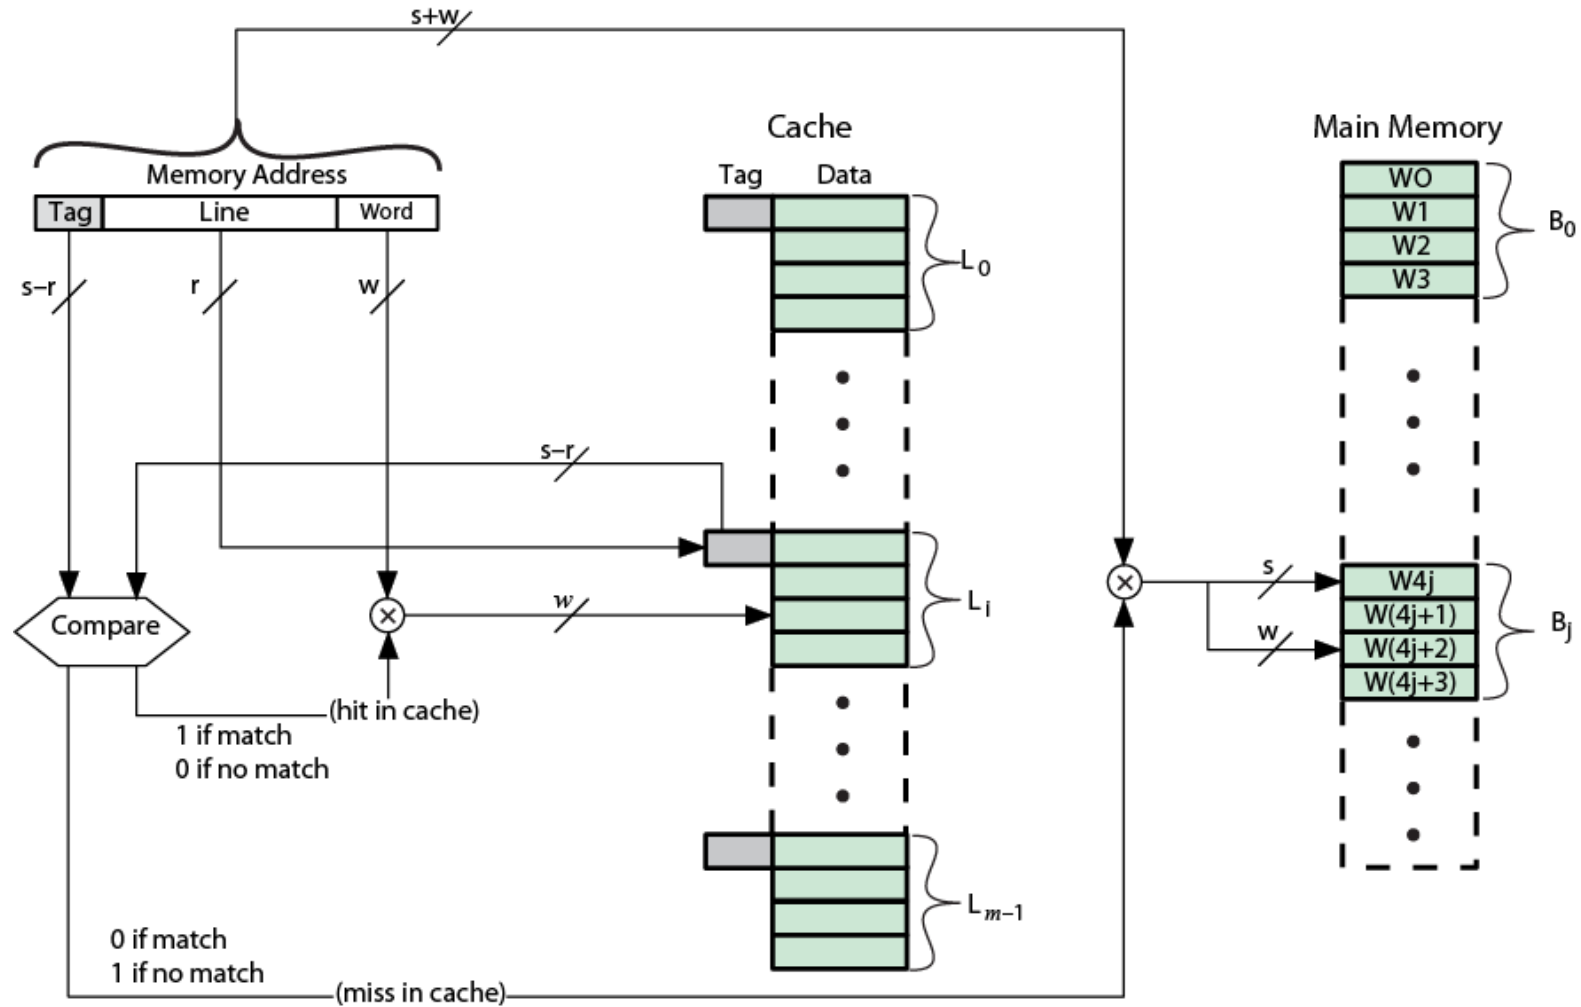


Main Memory

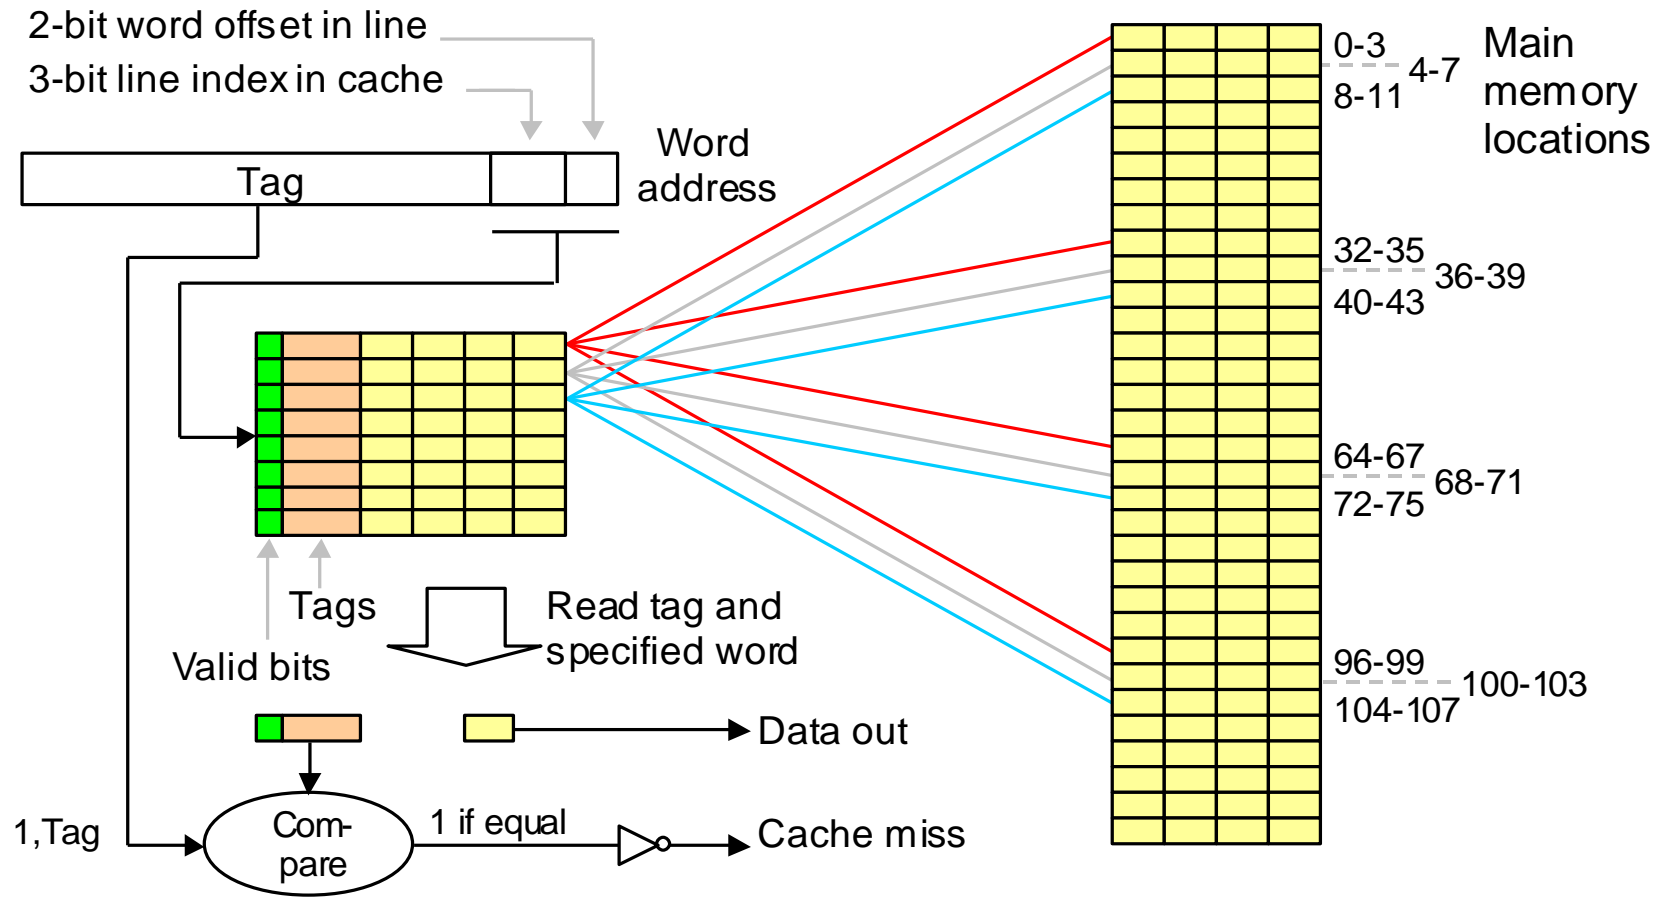


Memory Address

Direct Mapped Cache



Direct-Mapped Cache




Direct-mapped cache holding 32 words within eight 4-word lines. Each line is associated with a tag and a valid bit.

Example: Accessing a Direct-Mapped Cache

Show cache addressing for a byte-addressable memory with 32-bit addresses. Cache line $W = 16$ B. Cache size $L = 4096$ lines (64 KB).

Solution

Byte offset in line is $\log_2 16 = 4$ b. Cache line index is $\log_2 4096 = 12$ b. This leaves $32 - 12 - 4 = 16$ b for the tag.



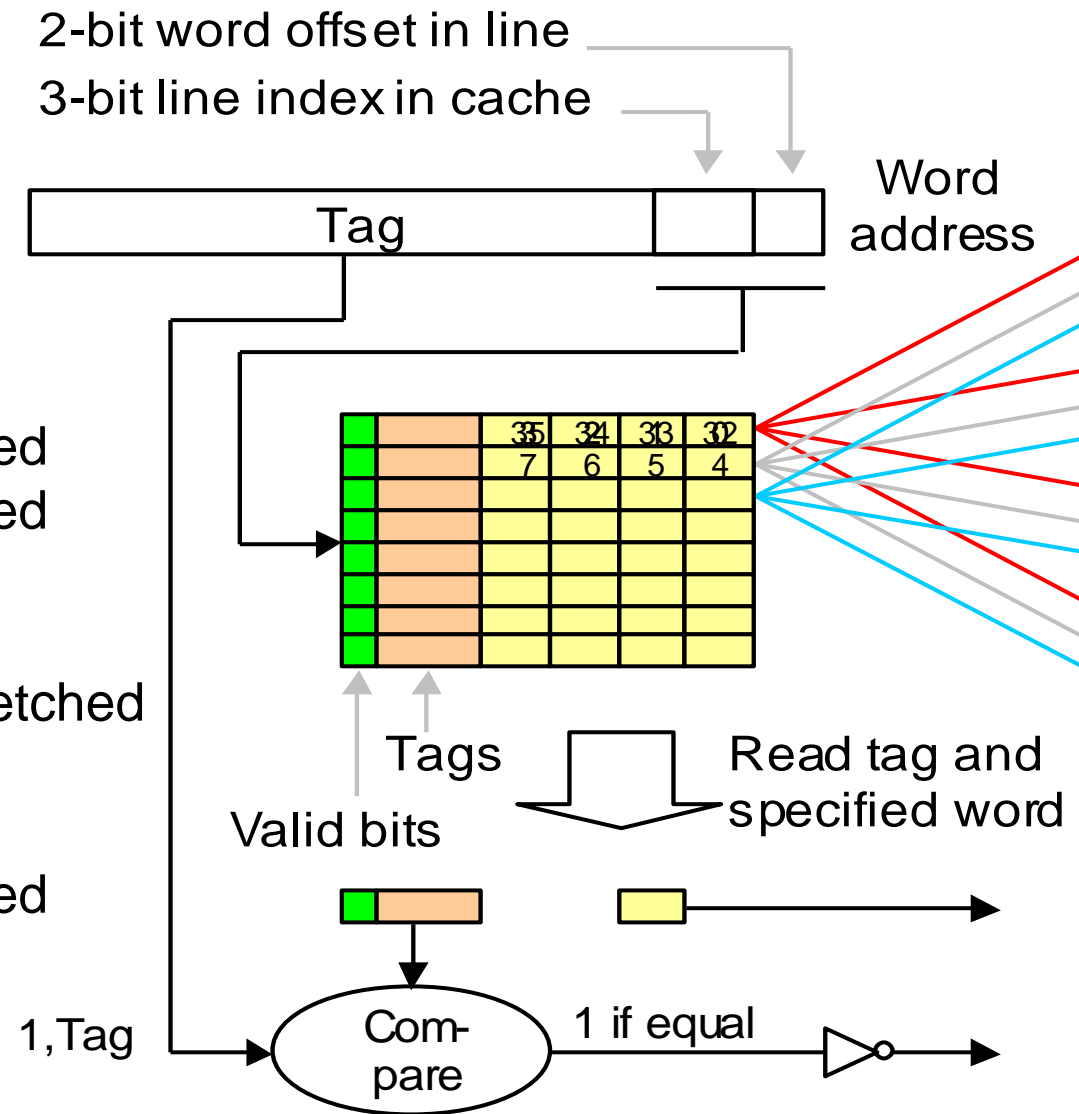
Components of the 32-bit address in an example direct-mapped cache with byte addressing.

Direct-Mapped Cache Behavior

Address trace:

1, 7, 6, 5, 32, 33, 1, 2, ...

1: miss, line 3, 2, 1, 0 fetched
7: miss, line 7, 6, 5, 4 fetched
6: hit
5: hit
32: miss, line 35, 34, 33, 32 fetched
(replaces 3, 2, 1, 0)
33: hit
1: miss, line 3, 2, 1, 0 fetched
(replaces 35, 34, 33, 32)
2: hit
... and so on



- Block replacement algorithm is trivial, as there is no choice.
- More than one MM block is mapped onto the same cache block.
 - May lead to contention even if the cache is not full.
 - New block will replace the old block.
 - May lead to poor performance if both the blocks are frequently used.
- The MM address is divided into three fields: *TAG*, *BLOCK* and *WORD*.
 - When a new block is loaded into the cache, the 8-bit *BLOCK* field determines the cache block where it is to be stored.
 - The high-order 11 bits are stored in a *TAG* register associated with the cache block.
 - When accessing a memory word, the corresponding TAG field is compared.
 - Match implies HIT.





(b) Associative Mapping

- Here, a MM block can potentially reside in *any cache block position*.
- The memory address is divided into two fields: **TAG** and **WORD**.
 - When a block is loaded into the cache from MM, the higher order 19 bits of the address are stored into the TAG register corresponding to the cache block.
 - When accessing memory, the 19-bit TAG field of the address is compared with all the TAG registers corresponding to all the cache blocks.
- Requires associative memory for storing the TAG values.
 - High cost / lack of scalability.
- Because of complete freedom in block positioning, a wide range of replacement algorithms is possible.

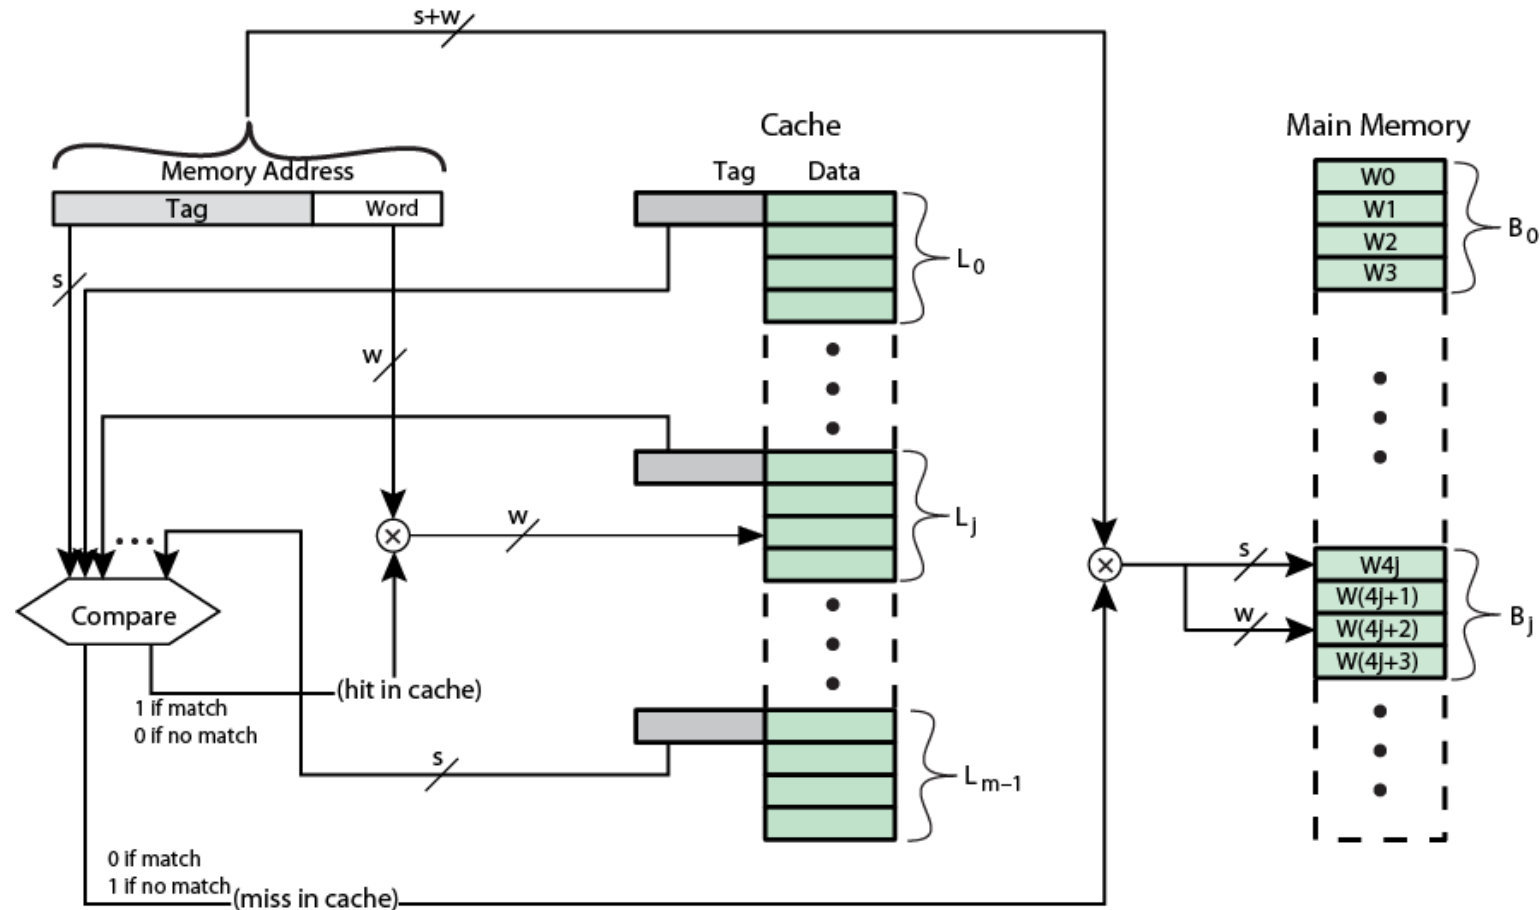
What is an Associative Memory?

- It is a memory where data is accessed by contents rather than address.
 - There is circuitry to compare the applied data value with all the stored values in memory in parallel.
 - Wherever there is a match, the memory system will be returning the information.
 - Very expensive to build – only small capacity units are feasible.

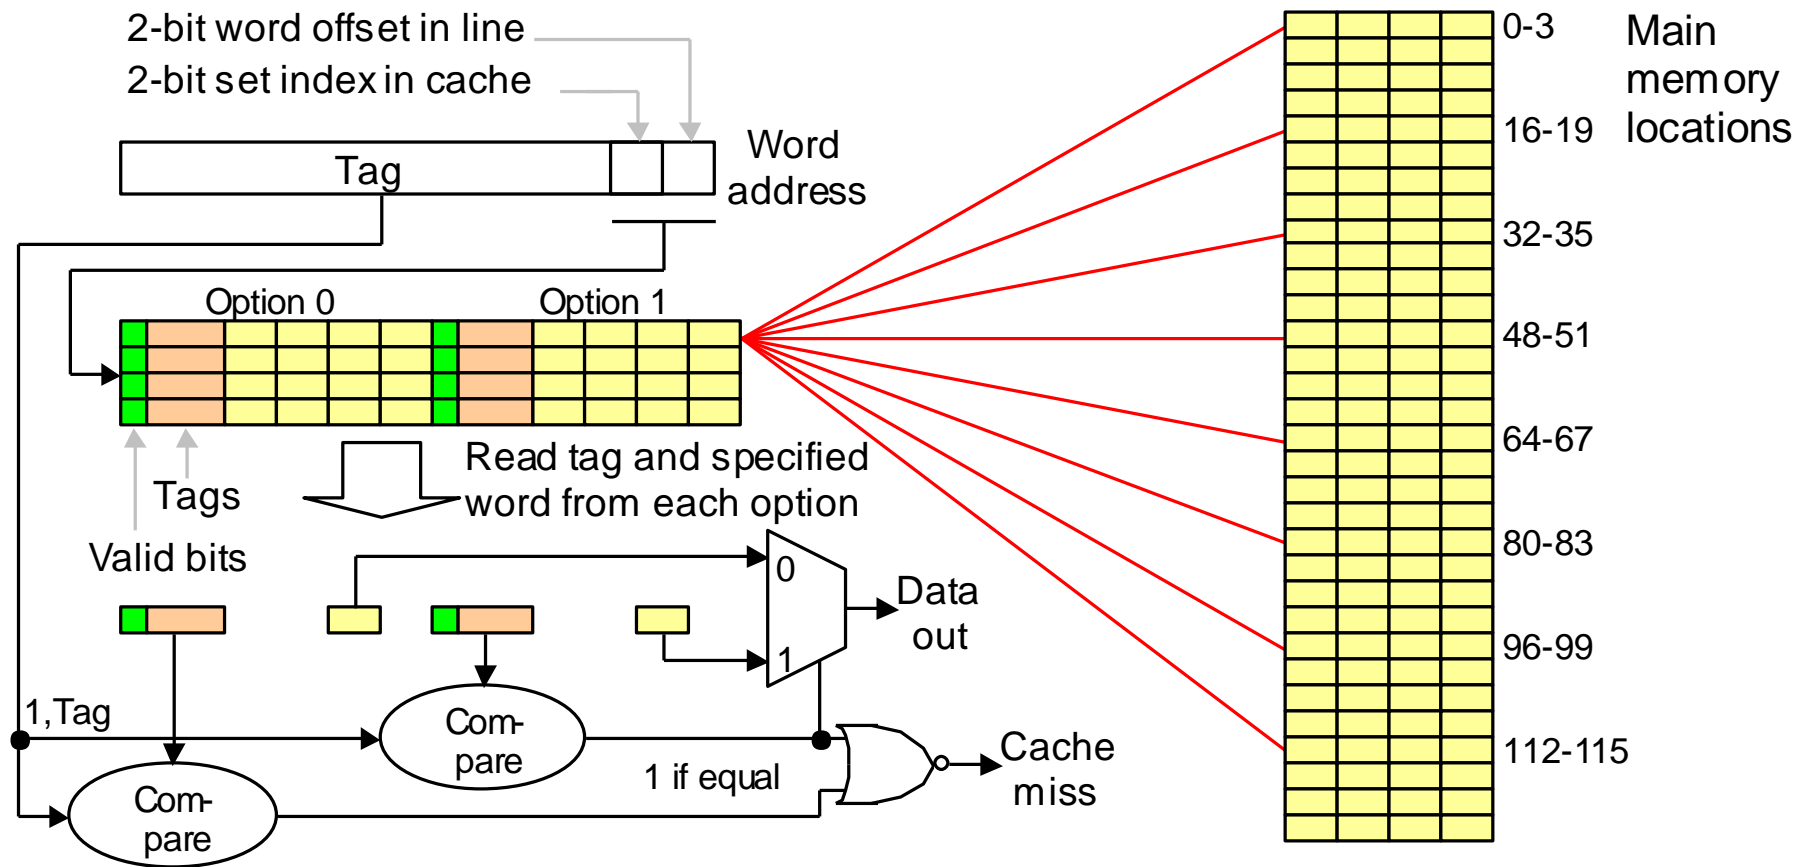
Fully Associative Mapping



Fully Associative Cache Design



Set-Associative Cache



Two-way set-associative cache holding 32 words of data within 4-word lines and 2-line sets.


Example: Accessing a Set-Associative Cache

Show cache addressing scheme for a byte-addressable memory with 32-bit addresses. Cache line width $2^W = 16$ B. Set size $2^S = 2$ lines. Cache size $2^L = 4096$ lines (64 KB).

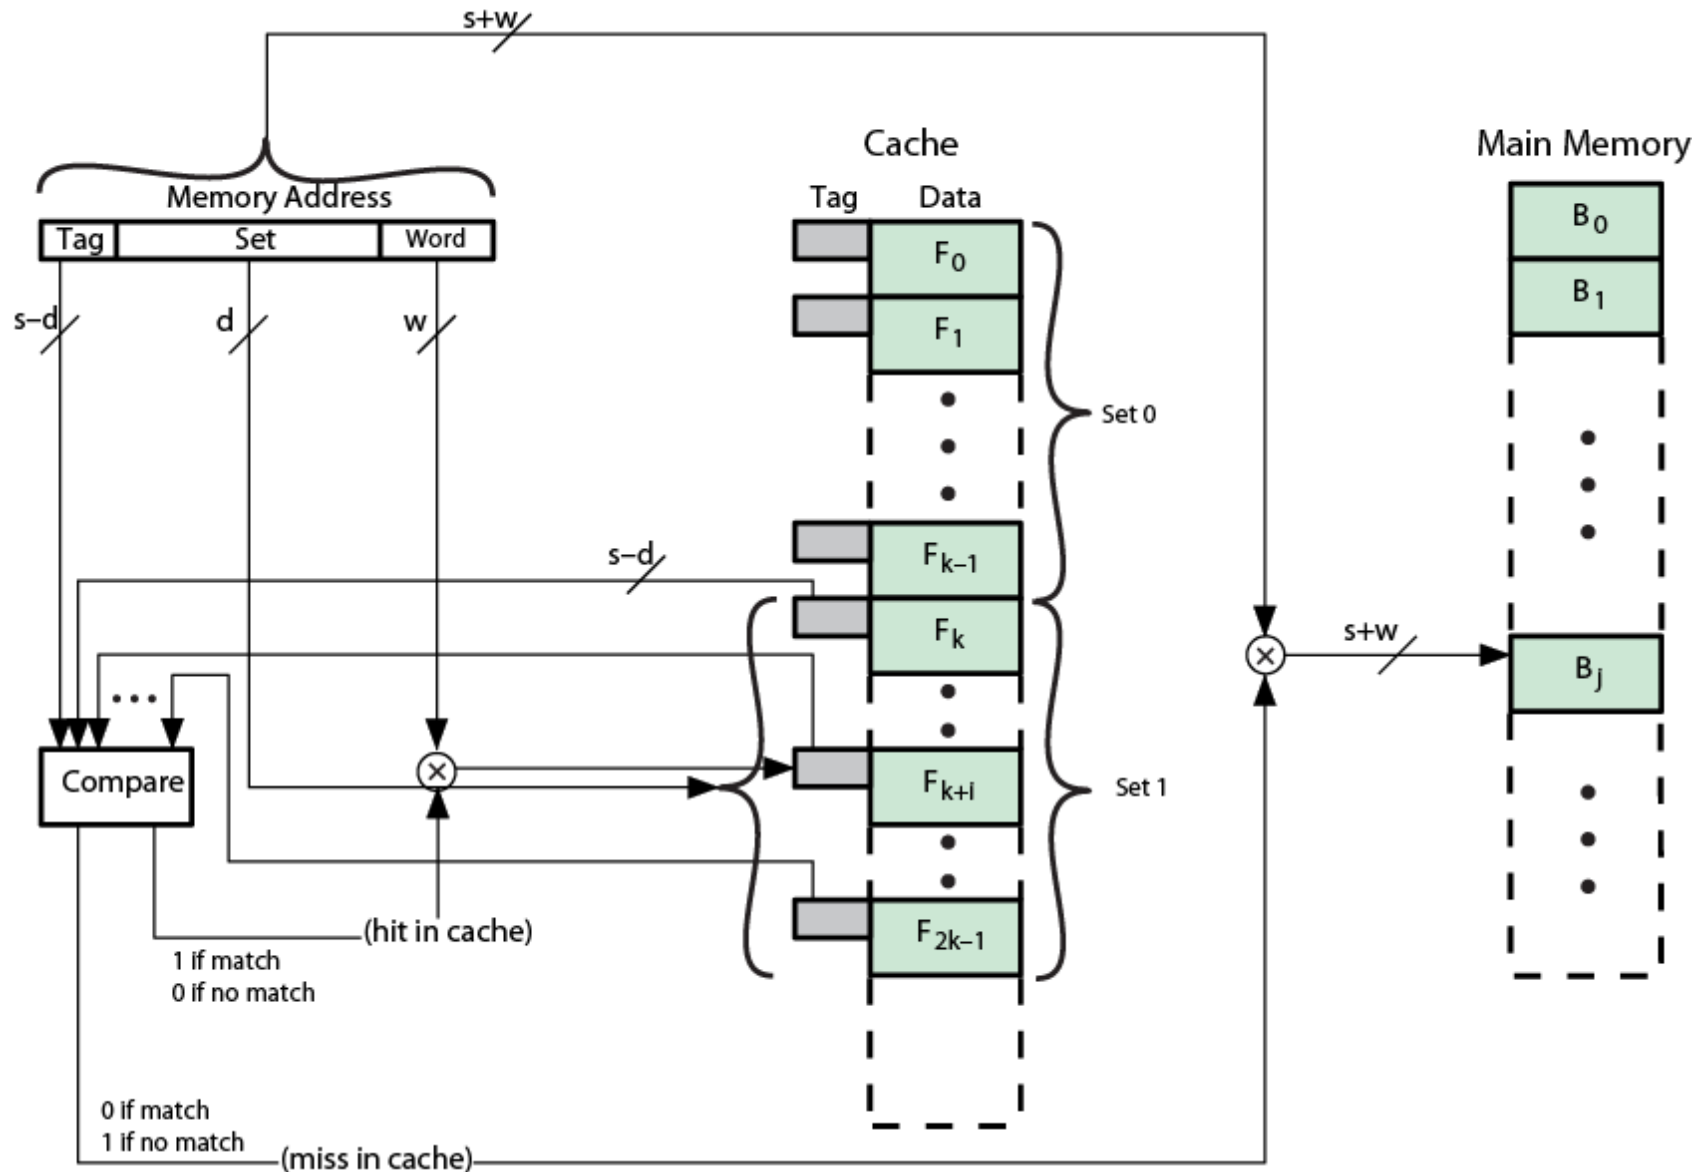
Solution

Byte offset in line is $\log_2 16 = 4$ b. Cache set index is $(\log_2 4096 / 2) = 11$ b. This leaves $32 - 11 - 4 = 17$ b for the tag.

Components of the 32-bit address in an example two-way set-associative cache.



K-way set-associative cache



Quiz 2: Cache Address Mapping


A 64 KB four-way set-associative cache is byte-addressable and contains 32 B lines. Memory addresses are 32 b wide.

- a. How wide are the tags in this cache?
- b. Which main memory addresses are mapped to set number 5?

Solution

The number of sets in the cache = $64\text{KB} / (4 \times 32\text{B}) = 512$.

- a. Address (32 b) = 5 b byte offset + 9 b set index + 18 b tag
- b. Addresses that have their 9-bit set index equal to 5. These are of the general form $2^{14}a + 2^5 \times 5 + b$; e.g., 160-191, 16 554-16 575, . .



Tag width =
 $32 - 5 - 9 = 18$


Set size = $4 \times 32 \text{ B} = 128 \text{ B}$
Number of sets = $2^{16} / 2^7 = 2^9$

Line width =
 $32 \text{ B} = 2^5 \text{ B}$



(c) N-way Set Associative Mapping

- A group of N consecutive blocks in the cache is called a *set*.
- This algorithm is a balance of direct mapping and associative mapping.
 - Like direct mapping, a MM block is mapped to a set.
$$\text{Set Number} = (\text{MM Block Number}) \% (\text{Number of Sets in Cache})$$
 - The block can be placed anywhere within the set (there are N choices)
- The value of N is a design parameter:
 - $N = 1$:: same as direct mapping.
 - $N = \text{number of cache blocks}$:: same as associative mapping.
 - Typical values of N used in practice are: 2, 4 or 8.



4-way Set Associative Mapping



Memory Address



Cache Memory



Main Memory

- Illustration for $N = 4$:
 - Number of sets in cache memory = 64.
 - Memory blocks are mapped to a set using modulo-64 operation.
 - Example: MM blocks 0, 64, 128, etc. all map to set 0, where they can occupy any of the four available positions.
- MM address is divided into three fields: *TAG*, *SET* and *WORD*.
 - The TAG field of the address must be associatively compared to the TAG fields of the 4 blocks of the selected set.
 - This instead of requiring a single large associative memory, we need a number of very small associative memories only one of which will be used at a time.

Q2. How is a block found if present in cache?

- Caches include a TAG associated with each cache block.
 - The TAG of every cache block where the block being requested may be present needs to be compared with the TAG field of the MM address.
 - All the possible tags are compared in parallel, as speed is important.
- Mapping Algorithms?
 - Direct mapping requires a single comparison.
 - Associative mapping requires a full associative search over all the TAGs corresponding to all cache blocks.
 - Set associative mapping requires a limited associated search over the TAGs of only the selected set.

- Use of *valid bit*:
 - There must be a way to know whether a cache block contains valid or garbage information.
 - A valid bit can be added to the TAG, which indicates whether the block contains valid data.
 - If the valid bit is not set, there is no need to match the corresponding TAG.

Q3. Which block should be replaced on a cache miss?

- With fully associative or set associative mapping, there can be several blocks to choose from for replacement when a miss occurs.
- Two primary strategies are used:
 - a) Random*: The candidate block is selected randomly for replacement. This simple strategy tends to spread allocation uniformly.
 - b) Least Recently Used (LRU)*: The block replaced is the one that has not been used for the longest period of time.
 - Makes use of a corollary of temporal locality:

“If recently used blocks are likely to be used again, then the best candidate for replacement is the least recently used block”

- To implement the LRU algorithm, the cache controller must track the LRU block as the computation proceeds.
- Example: Consider a 4-way set associative cache.
 - For tracking the LRU block within a set, we use a 2-bit counter with every block.
 - When hit occurs:
 - Counter of the referenced block is reset to 0.
 - Counters with values originally lower than the referenced one are incremented by 1, and all others remain unchanged.
 - When miss occurs:
 - If the set is not full, the counter associated with the new block loaded is set to 0, and all other counters are incremented by 1.
 - If the set is full, the block with counter value 3 is removed, the new block put in its place, and the counter set to 0. The other three counters are incremented by 1.

- It may be verified that the counter values of occupied blocks are all distinct.
- An example:

X	Block 0
X	Block 1
X	Block 2
X	Block 3

Initial

X	Block 0
X	Block 1
0	Block 2
X	Block 3

Miss: Block 2

0	Block 0
X	Block 1
1	Block 2
X	Block 3

Miss: Block 0

1	Block 0
X	Block 1
2	Block 2
0	Block 3

Miss: Block 3

2	Block 0
0	Block 1
3	Block 2
1	Block 3

Miss: Block 1

0	Block 0
1	Block 1
3	Block 2
2	Block 3

Hit: Block 0

1	Block 0
2	Block 1
0	Block 2
3	Block 3

Miss: Block 2

2	Block 0
3	Block 1
1	Block 2
0	Block 3

Hit: Block 3

2	Block 0
3	Block 1
0	Block 2
1	Block 3

Hit: Block 2

0	Block 0
3	Block 1
1	Block 2
2	Block 3

Hit: Block 0

1	Block 0
0	Block 1
2	Block 2
3	Block 3

Miss: Block 1

1	Block 0
0	Block 1
2	Block 2
3	Block 3

Hit: Block 1

Types of Cache Misses

a) **Compulsory Miss**

- On the first access to a block, the block must be brought into the cache.
- Also known as cold start misses, or first reference misses.
- Can be reduced by increasing cache block size or prefetching cache blocks.

b) **Capacity Miss**

- Blocks may be replaced from cache because the cache cannot hold all the blocks needed by a program.
- Can be reduced by increasing the total cache size.

c) **Conflict Miss**

- In case of direct mapping or N-way set associative mapping, several blocks may be mapped to the same block or set in the cache.
- May result in block replacements and hence access misses, even though all the cache blocks may not be occupied..
- Can be reduced by increasing the value of N (cache associativity).

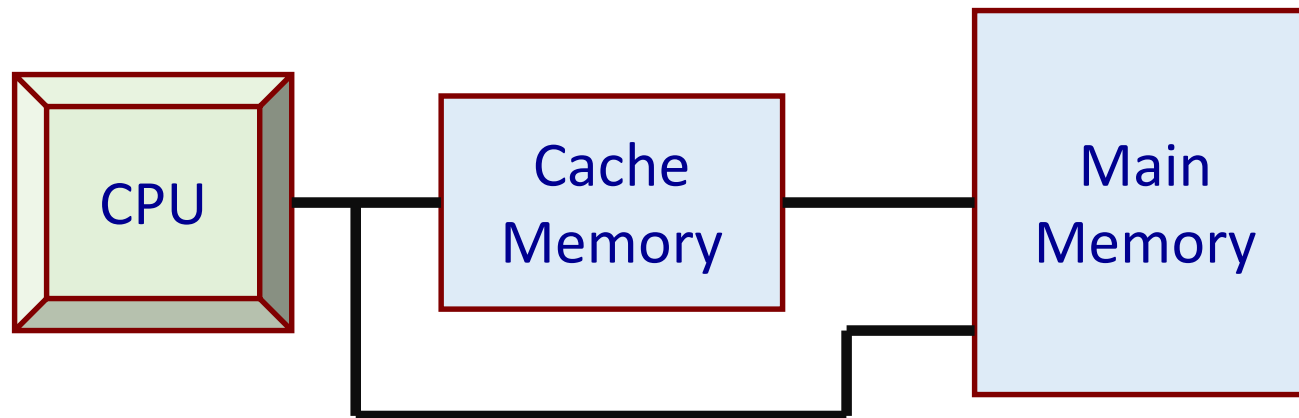
Q4. What happens on a write?

- Statistical data suggests that read operations (including instruction fetches) dominate processor cache accesses.
 - All instruction fetch operations are read.
 - Most instructions do not write to memory.
- Making the common case fast:
 - Optimize cache accesses for reads.
 - But Amdahl's law reminds that for high performance designs we cannot ignore the speed of write operations.

- The common case (read operations) is relatively easy to make faster.
 - A block(s) can be read at the same time while the TAG is being compared with the block address.
 - If the read is a HIT the data can be passed to the CPU; if it is a MISS ignore it.
- Problems with write operations:
 - The CPU specifies the size of the write (between 1 and 8 bytes), and only that portion of a block has to be changed.
 - Implies a read-modify-write sequence of operations on the block.
 - Also, the process of modifying the block cannot begin until the TAG is checked to see if it is a hit.
 - Thus, cache write operations take more time than cache read operations.

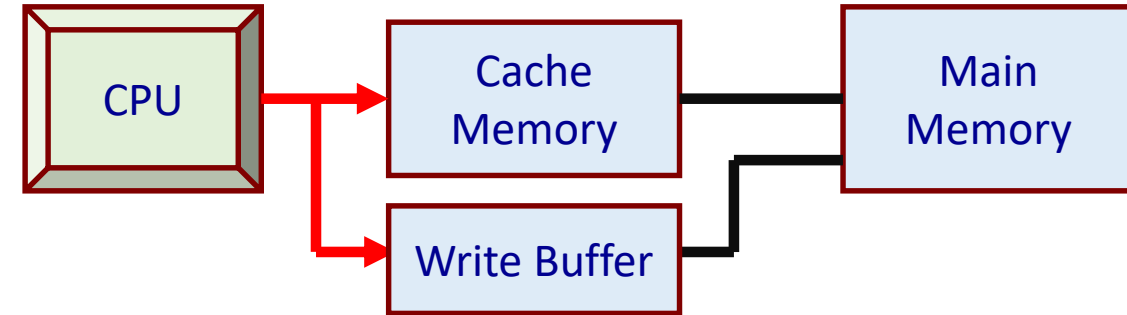
Cache Write Strategies

- Cache designs can be classified based on the write and memory update strategy being used.
 1. Write Through / Store Through
 2. Write Back / Copy Back



(a) Write Through Strategy

- Information is written to both the cache block and the main memory block.
- Features:
 - Easier to implement
 - Read misses do not result in writes to the lower level (i.e. MM).
 - The lower level (i.e. MM) has the most updated version of the data – important for I/O operations and multiprocessor systems.
 - A write buffer is often used to reduce CPU write stall time while data is written to main memory.



- **Perfect Write Buffer:**

- All writes are handled by write buffer; no stalling for write operations.
- For unified L1 cache.

$$\text{Stall Cycles / Memory Access} = \% \text{ Reads} \times (1 - H_{L1}) \cdot t_{MM}$$

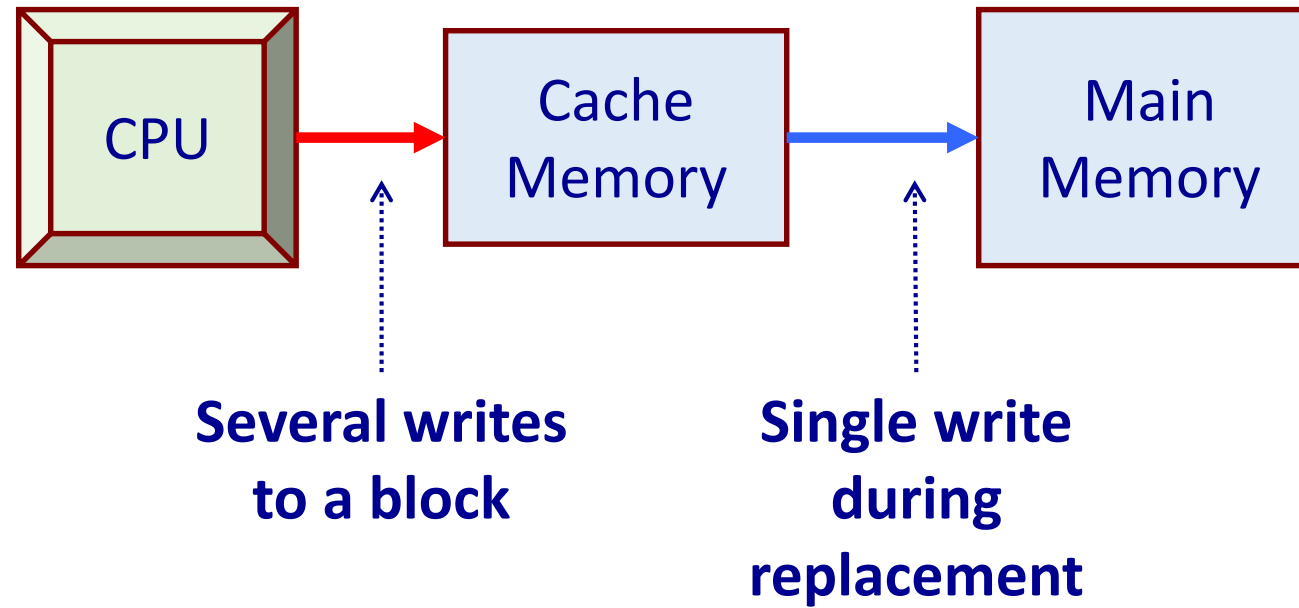
- **Realistic Write Buffer:**

- A percentage of write stalls are not eliminated when the write buffer is full.
- For unified L1 cache,

$$\text{Stall Cycles / Memory Access} = (\% \text{ Reads} \times (1 - H_{L1}) + \% \text{ write stalls not eliminated}) \times t_{MM}$$

(b) Write Back Strategy

- Information is written only to the cache block.
- A modified cache block is written to MM only when it is replaced.
- Features:
 - Writes occur at the speed of cache memory.
 - Multiple writes to a cache block requires only one write to MM.
 - Uses less memory bandwidth, makes it attractive to multiprocessors.
- Write-back cache blocks can be *clean* or *dirty*.
 - A status bit called *dirty bit* or *modified bit* is associated with each cache block, which indicates whether the block was modified in the cache (0: clean, 1: dirty).
 - If the status is clean, the block is not written back to MM while being replaced.



Cache Write Miss Policy

- Since information is usually not needed immediately on a write miss, two options are possible on a cache write miss:

a) Write Allocate

- The missed block is loaded into cache on a write miss, followed by write hit actions.
- Requires a cache block to be *allocated* for the block to be written into.

b) No-Write Allocate

- The block is modified only in the lower level (i.e. MM), and not loaded into cache.
- Cache block is *not allocated* for the block to be written into.

- Typical usage:
 - a) Write-back cache with write-allocate
 - In order to capture subsequent writes to the block in cache.
 - b) Write-through cache with no-write-allocate
 - Since subsequent writes still have to go to MM.

Estimation of Miss Penalties

- **Write-Through Cache**

- Write Hit Operation:
 - Without write buffer, miss penalty = t_{MM}
 - With perfect write buffer, miss penalty = 0

- **Write-Back Cache**

- Write Hit Operation
 - Miss penalty = 0

- **Write-Back Cache (with Write Allocate)**

- **Write Hit Operation**

- Miss penalty = 0

- **Read or Write Miss Operation**

- If the replaced block is clean, miss penalty = t_{MM}
 - No need to write the block back to MM.
 - New block to be brought into MM (t_{MM}).
 - If the replaced block is dirty, miss penalty = $2 t_{MM}$
 - Write the block to be replaced to MM (t_{MM}).
 - New block to be brought into MM (t_{MM}).