

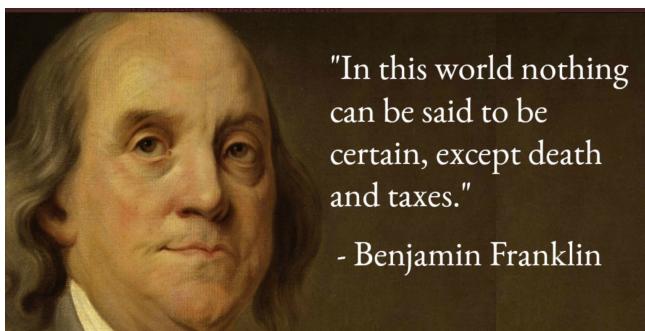
## • Algorithms under Uncertainty .

- What makes life interesting ?



Decision-making  
under uncertainty .

Uncertainty is  
omnipresent !



- Whom to marry ?
- Which classes to take ?
- Should I buy or rent ?
- Which stocks to buy ?

### Algorithms under uncertainty :

- Decisionmaker don't have complete knowledge of the input.
- Each time step, decisionmaker needs to take a decision based only on knowledge of the past .

### Several approaches :

- Online Algorithms.
- Online Learning.
- Online convex optimization.

Prerequisite :

Basics of prob.  
theory & linear  
programming .

- **Online learning**: Answer a sequence of questions given (partial) knowledge of the correct answers to previous questions and possibly additional available information.

## Online Learning

for  $t = 1, 2, \dots$

receive question  $\mathbf{x}_t \in \mathcal{X}$  instance domain

predict  $p_t \in D \subseteq \mathcal{Y}$

receive true answer  $y_t \in \mathcal{Y}$  target domain

suffer loss  $l(p_t, y_t)$

$$\begin{array}{lll} \mathbf{x}_1 & p_1 & y_1 \\ \mathbf{x}_2 & p_2 & y_2 \\ \vdots & & \end{array}$$

Goal: Minimize cumulative loss (regret) suffered along its run.

Example ( Predicting whether it is going to rain tomorrow: )

day  $t$ , the question  $x_t$  can be encoded as a vector of meteorological measurements

the learner should predict if it's going to rain tomorrow output a prediction

in  $[0, 1]$ ,  $D \neq \mathcal{Y}$   $\{0, 1\}$

loss function:  $\ell(p_t, y_t) = |p_t - y_t|$

$$\begin{array}{cc} 0.7 & 0 \\ 0.6 & 1 \end{array} \quad \boxed{\begin{array}{c} 0.7 \\ + \\ 0.4 \end{array}}$$

which can be interpreted as the probability to err if predicting that it's going to rain with probability  $p_t$

## • Online Convex Optimization (OCO) :

Online Convex Optimization (OCO)

**input:** A convex set  $S$

**for**  $t = 1, 2, \dots$

predict a vector  $\mathbf{w}_t \in S$

receive a convex loss function  $f_t : S \rightarrow \mathbb{R}$

suffer loss  $f_t(\mathbf{w}_t)$

$$\text{Regret}_{\tau}(\mathbf{u}) = \sum_{t=1}^{\tau} f_t(\mathbf{w}_t) - \sum_{t=1}^{\tau} f_t(\mathbf{u})$$

*Best hypothesis*

Example ( Prediction from expert advice )

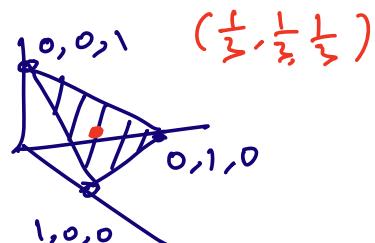
The decision maker has to choose among the advice of  $n$  given experts. i.e., the  $n$ -dimensional simplex  $\mathcal{X} = \{x \in \mathbb{R}^n, \sum_i x_i = 1, x_i \geq 0\}$ .

$g_t(i)$ : the cost of the  $i$ 'th expert at iteration  $t$

$g_t$ : the cost vector of all  $n$  experts

The cost function is given by the linear function  $f_t(w) = g_t^T x$ .

There are  $n$  stocks .



- This lecture & next few weeks, we focus on Online Algorithms.
- Decisionmaker don't have complete knowledge of the input.
- Input arrives over time : part by part.
- Each time step, decision maker needs to take a decision : immediate & irrevocable, based only on knowledge of the past.

- Performance Measure :

Competitive Ratio.

An online algorithm  $\text{ALG}$  is  $c$ -competitive if there is a constant  $c$  s.t. for all finite input sequences  $I$ ,

$$\text{ALG}(I) \leq c \cdot \text{OPT}(I) + \alpha.$$

- For additive constant  $\alpha=0$ , we call  $\text{ALG}$  is strictly  $c$ -competitive.

- Mostly we are interested in efficient (polytime) algorithms.
- But hardness also lies in the information theoretic barrier!

- Paging Problem:

Consider two-level computer memory system.



- At each time step  $t$  a request for some page  $P_{it}$  comes.
  - A **hit** occurs if  $P_{it}$  is already in cache.
  - Otherwise, a **miss** occurs ;  
Then the system incurs one page fault and  $P_{it}$  must be fetched from slow memory to cache.
  - Goal: Decide which  $k$  pages to retain in the cache, at each point of time, to minimize misses / maximize hits.
  - Typically, for paging
    - fast memory = RAM,
    - slow memory = Disk.
  - In Caching (where page is called Block),
    - fast memory = Cache,
    - slow memory = RAM.
- We use caching/paging interchangeably.

- Multiple abstract cost model :  $f \ll s$   
 General: Cost  $f$  for reading fast memory,  
 Cost  $s$  for fetching from slow memory.
- Page fault model :  $f = 0, s = 1$ . — Our model!
- Full access model :  $f = 1, s = s$ .

### Candidate algorithms :

- Online
- LRU (LEAST-RECENTLY-USED): When eviction is necessary, replace the page whose most recent request was earliest.
  - CLOCK (CLOCK-REPLACEMENT): An approximation to LRU in which a single “use bit” replaces the implicit (time of last access) timestamp of LRU.<sup>1</sup>
  - FIFO (FIRST-IN/FIRST-OUT): Replace the page that has been in the fast memory longest.
  - LIFO (LAST-IN/FIRST-OUT): Replace the page most recently moved to the fast memory.
  - LFU (LEAST-FREQUENTLY-USED): Replace the page that has been requested the least since entering the fast memory.
  - LFD (LONGEST-FORWARD-DISTANCE): Replace the page whose next request is latest.

Offline

Assume  $K = 3$ .

Requests :  $P_4, P_2, P_1, P_4, P_3, P_2, P_1, P_4$ .

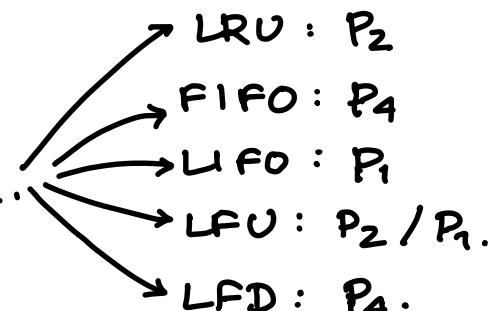
At  $t = 3$ .

$P_4$	$P_2$	$P_1$
-------	-------	-------

$t = 4$ .

$P_4$	$P_2$	$P_1$
-------	-------	-------

$t = 5$ , Need to evict a page.



Note: All are

Demand paging algorithms, i.e., unless there is a page fault they never evict a page from the cache.

HW: Prove that any online/offline paging algorithm can be modified to be demand paging without increasing the overall cost on any request sequence.

Corollary 5 There is a paging algorithm such that

Daniely-Mansour '19

- Its competitive ratio is  $2H_k$  → compares w. best offline algo
- Its regret on any interval  $I$  is → compares w. single fixed cache

$$O\left(k\sqrt{|I|\log(T)} + \sqrt{k|I|\log\left(\binom{n}{k}T\right)}\right) = O\left(k\sqrt{|I|\log(nT)}\right)$$

Let us rephrase Corollary 5 in terms of the paging problem. It guarantees an online algorithm for which the number of cache misses is at most a  $2H_k$  larger than what is achieved by the optimal offline schedule. Likewise, for long enough segments, i.e., longer than  $\Omega(k^2 \log(nT))$ , it guarantees that the number of cache misses is not much larger compared to the best single “fixed” cache  $C^*$ . Recall that for a fixed cache  $C^*$ , whenever we have a request for a page  $i \notin C^*$  we first fetch  $i$  evicting  $j \in C^*$  and then evict  $i$  and fetch back  $j$ , for a total cost of two.

## § Optimal offline paging algorithm:

LFD.

Intuition: Any optimal offline paging algo  $A$  can be modified to act like LFD without degrading its performance.

Claim: Let  $\sigma$  be any request sequence.

$\forall i \in [1|\sigma|]$ , we can construct  $A_i$  s.t. → offline algo

- (i)  $A_i$  processes first  $i-1$  requests same as  $A$ .
- (ii) If  $i$ 'th request results in a page fault,  $A_i$  evicts from its fast memory the page with LFD.
- (iii)  $A_i(\sigma) \leq A(\sigma)$ .

Claim implies LFD is optimal.

Start with  $A = OPT \xrightarrow{\text{claim } i=1} OPT_1 \xrightarrow{\text{claim } i=2} \dots \xrightarrow{} OPT_n = LFD$ .

Proof of Claim:

Assume just after processing  $i$ 'th request, fast memories of  $A$  &  $A_i$  contains page sets

$X \cup \{v\}$ ,  $X \cup \{u\}$  resp.



$A_i$  has evicted  $v$   
 $A$  " "  $u$

W.l.o.g. assume  $v \neq u$  ( $i$ th request results in page-fault)

Until  $v$  is requested, for subsequent requests.

$A_i$  mimics  $A$  except for evicting  $u$  if  $A$  evicts  $v$ .

Note: # common pages is always  $\geq k-1$ .

If # common pages becomes  $k$  (i.e., if  $A$  evicts  $v$ ) then  $A_i$  continues same as  $A$ .

However, if  $v$  is requested before  $A$  evicts  $v$ , then  $A_i$  incurs page fault, but not  $A$ .

However, when  $v$  was evicted by  $A_i$  it must have LFD, so  $\exists$  at least one request for  $u$  after the  $i$ 'th page fault.

That incurs a page fault to  $A$  but not to  $A_i$ .

$\therefore$  # page faults for  $A_i \leq$  # page faults for  $A$ .  
(after servicing  $v$ )

& after servicing  $v$ ,  $A$  and  $A_i$  identify.  $\square$

- Polytime solvable!

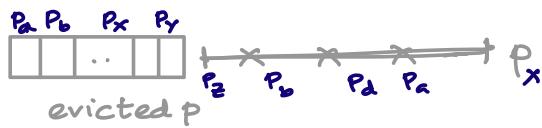
- Hardness comes from lack of information.

## § Lower Bound of $K$ for deterministic algorithms for paging.

Claim : For any finite sequence  $\sigma$  chosen from a set of  $K+1$  pages,  $LFD(\sigma) \leq |\sigma|/K$ .

→ LFD faults at most once every  $K$  requests.

[ If LFD evicts a page  $p_x$ , all other pages in cache must be requested prior to next request for  $p_x$  ]



Pf of lower bound : pages :  $p_1, p_2, \dots, p_{K+1}$ .

w.l.o.g. assume an online <sup>(deterministic)</sup> algo A holds  $p_1, \dots, p_K$ .

We define request sequence  $\sigma$  inductively :

$$r_1 = p_{K+1},$$

$r_{i+1}$  = unique page not in A's cache just after servicing  $r_1, \dots, r_i$ .



Note: A faults on each request & we can make  $|\sigma| \rightarrow \infty$ .

$$\therefore A(\sigma) = |\sigma|, \quad OPT(\sigma) = LFD(\sigma) \leq |\sigma|/K.$$

$$\Rightarrow C.R.(A) \geq \frac{|\sigma|}{|\sigma|/K} = K.$$

- Paging is not hard due to computational reasons, but due to information theoretic reasons.

- The  $(h, k)$ -paging problem:  $h, k \in \mathbb{Z}, h \leq k$ .

Measure performance of online paging algo of cache size =  $k$ , relative to optimal offline algo with cache  $h \leq k$ .

↙ weaker adversary!

Belady's anomaly: Interestingly, on some input sequence some algorithms (e.g. FIFO) may perform better when it has smaller fast memory.

(HW: On properties of different algorithm)

### § Marking Algorithms:

- K-phase partition: Fast memory size =  $k$   
Divide  $\sigma$  into phases. Phase 0 = empty sequence.  
Phase  $i$  is the maximal sequence following phase  $i-1$  that contains  $\leq k$  distinct page requests.  
So phase  $i+1$  begins with  $(k+1)$ th distinct pg rq.

$$K=2 \quad | P_1 P_3 P_1 | P_2 P_3 P_2 P_3 P_2 | P_4 P_1 P_2 P_3$$

- Associate a bit "mark" for each page.  
Bit is set → Marked, else → Unmarked.
- Beginning of each K-phase:  
unmark all pages currently in fast memory.
- During a K-phase,  
Mark a page when it is first requested during the phase.
- A marking algorithm never evicts a marked page from its fast memory!



- LRU, CLOCK, and many more algorithms belong to the class of marking algorithms.
- We'll show any marking algorithm attains optimal competitive ratio in the page-fault model.

Thm: LRU is marking algorithm.

- for contradiction, assume LRU evicts a marked page  $x$  during some  $k$ -phase.
- Consider the first request for  $x$  during this  $k$ -phase. Immediately after serving  $x$ : it is marked as most recently used page.
- So to evict  $x$  by LRU, there should be  $\geq k+1$  different pages to be requested in this phase.  
[ $k$  pages including  $x$  + the page which got  $x$  evicted]



we need  $k$  other pages to be more recently used.

- contradiction with defn of phase!

{ Theorem : Any marking algo ALG is  $\frac{K}{K-h+1}$  competitive.

→ Any request sequence  $\sigma$  & its K-phase partition.

Claim : For any phase  $i$ , ALG incurs  $\leq K$  pg faults

→ There are  $K$  distinct pg references in each phase. Once a pg is accessed, it is marked & can't be evicted till end of phase.

so, ALG can't fault twice on same page.

Let  $q$  be first request of phase  $i$ , consider sequence  $\sigma'$  starting w. second req. of phase  $i$  upto and including first req. of phase  $(i+1)$ .

OPT has  $h-1$  pages excluding  $q$ .

There are  $K$  distinct requests in  $\sigma'$ .

So, OPT must incur  $\geq K-(h-1) = K-h+1$  faults.

(We can ignore last phase by additive  $O(K)$  term?)

In each phase alg incurs fault  $\leq K$  times.

$$ALG(\sigma) \leq \frac{K}{K-h+1} \cdot OPT(\sigma) + K.$$

upper bound from last phase. ■

## Power of Randomization:

- Adversary Models.

'Natural' model in deterministic case:

- knows the algo & chooses worst-case input to maximize C.R.

- Advantage of randomness : Adversary may not know the outcome of the random choice.

- Adversary constructs a sequence  $\sigma$  <sup>request</sup> & pays a cost.

### ① OBL (oblivious) : Weak Adversary

- constructs  $\sigma$  in advance.
- pays optimally.
- Doesn't know actions (due to outcomes of random choices) by Algo.

### ② ADON (adaptive-online) : Medium Adversary

- Generate  $\sigma(t)$  based on Algo's action so far.  
(on  $\sigma(1) \dots \sigma(t-1)$ ).
- serves current request online.

[So it knows its own strategy for generating  $\sigma$ ,  
Description of online algo & its action so far;  
& then need to perform in online manner].

### ③ ADOF (adaptive-offline) : Strong Adversary

- Choose next request based on Algo's action so far.
- Pays optimal offline cost (on  $\sigma$  that is created online by adversary).

- Competitive ratio:  $\forall \sigma, \mathbb{E}[\text{ALG}(\sigma)] \leq c \cdot \text{OBL}(\sigma) + \alpha =: r(\text{OBL}(\sigma))$ .  
 OPT( $\sigma$ )  $\xrightarrow{\text{some constant}}_{\text{indep of } \sigma}$

Expectation is taken over the random choices made by ALG.

Note: ADV( $\sigma$ ) itself becomes a RV for ADON/ADOF. So a more careful definition is needed there.

- Adversary:  $(Q, S)$ .  
 $\nwarrow$  servicing component  
 (answer requests created by Q)  
 $\searrow$  requesting component  
 (creates request)

S for OBL / AD-OF : Optimal offline algo.  
 (so completely determined by Q).

Q for OBL : Fix sequence  $\sigma$  that may depend on the decision-makers online algo.

Q for AD-ON    Q is a sequence of functions  
 or AD-OF :  $q_i : \underbrace{A_1 \times A_2 \times \dots \times A_{i-1}}_{\text{Prev. actions}} \rightarrow R$ .

Let cost of ALG against ADV :  $\text{ALG}(\text{ADV}) = \text{ALG}(\sigma)$   
 where  $\sigma$  is constructed by Q via interaction w. ALG.

Similarly, adversary cost against ALG =  $\text{ADV}(\text{ALG})$ .

AD-OFF:  $\tau$ -comp against AD-OFF if  $\nexists$  AD-OFF  $Q$ :

$$\mathbb{E}[\text{ALG}(Q)] = \mathbb{E}_x \left[ \text{ALG}_x \left[ \sigma(\text{ALG}_x, Q) \right] \right]$$

$$\leq \tau \mathbb{E}_x \left[ \text{OPT} \left[ \sigma(\text{ALG}_x, Q) \right] \right] = \tau \mathbb{E} \left[ \text{OPT} \left[ \sigma(\text{ALG}, Q) \right] \right].$$

AD-ON:  $\tau$ -comp against AD-ON if  $\nexists$  AD-ON  $(Q, S)$

$$\mathbb{E}[\text{ALG}(\text{ADV})] = \mathbb{E}_x [\text{ALG}_x (\sigma(\text{ALG}_x, Q))]$$

$$\leq \tau (\mathbb{E}_x [\text{ADV}(\text{ALG}_x)]) = \tau (\mathbb{E} [\text{ADV}(\text{ALG})]).$$

### § Relating the adversaries: (Ch 7.3 in B-T)

Theorem: If there exists a randomized online algorithm that is  $c$ -competitive against adaptive offline adversaries, then there also exists a  $c$ -competitive deterministic online algorithm.

Theorem: If  $A$  is  $c$ -competitive against adaptive online adversaries and there exists a  $d$ -competitive algo against oblivious adversaries, then there exists a  $cd$ -competitive algorithm against adaptive offline adversaries.

Corollary: If  $A$  is  $c$ -competitive against adaptive online adversaries, then there exists a  $c^2$ -competitive deterministic algorithm.

- Randomized Algorithm:

- Algorithm RAND: Whenever a page fault occurs, evict a page chosen randomly and uniformly among all fast memory pages.

→ RAND can be shown to be  $\frac{k}{k-h+1}$ -competitive.

§ The optimal algorithm: MARK.

Never evicts  
a marked pg.

- Initially, all the pages are marked.
- If there is a request for a page  $p$ ,
  - If  $p$  in cache, but unmarked, then mark  $p$ .
  - Else if  $p$  is not in cache, then  $p$  is brought into cache, replacing a randomly and uniformly chosen page from the set of all unmarked pages. Mark  $p$ .

If all pages in cache are marked when  $p$  is about to be brought in, then they are all unmarked first.

Mark  $p$ .

↙ K-memory bit

- Thm: C.R. of MARK (against obl. adversaries) is  $O(\log k)$ .

Proof: W.l.o.g. assume MARK & adversary start with the same set of initial pages in the cache (else we have an additive  $O(k)$  term).

Fix a request sequence  $J$  & its  $k$ -phase partition.

for each phase  $i$ , at start pages in cache are called old pages.

Non-old page requests are new pages.

Obs: (i) Each request is either old/new pg.

(ii) Repeated requests don't contribute to online cost.

(iii) Beginning of phase, all old pages are unmarked.

→  $K$  distinct requests give  $K$  marked pages.

→ New pg not in cache, so all are unmarked.

Consider phase  $i$ . Has  $K$  distinct pg requests.

$m_i$  : # new pages requested in this phase.

Worst possible sequence  $J$  : Request all new pages ( $m_i$  page faults), followed by  $(K - m_i)$  (first) requests to old pages (that were evicted).

→ Adv doesn't know which pages are being evicted.

Key!  
Obs: In  $i$ 'th phase,  $j$ 'th requested old page is in cache w.p.  $= \frac{(K - m_i - (j-1))}{(K - (j-1))}$ .

# old unmarked pgs in cache      # old unmarked pages  
only  $j-1$  requested ones are marked

$$\text{Hence, } \Pr[\text{fault event}] = 1 - \frac{K - m_i - (j-1)}{K - (j-1)} = \frac{m_i}{K - j + 1}.$$

$$\begin{aligned} \mathbb{E}[\text{faults during } i\text{'th phase}] &= m_i + \sum_{j=1}^{K-m_i} \frac{m_i}{K-j+1} \\ &= m_i + m_i (H_K - H_{m_i}) \leq m_i H_K. \end{aligned}$$

$$\text{Here, } H_K = \sum_{i=1}^K \frac{1}{i}.$$

( HW: Show  $\ln K \leq H_K \leq 1 + \ln K$ .

Intuition:  $\int_1^K \frac{1}{x} dx \approx \ln K$ ).

Hence, MARK pays  $\leq \sum_{i=1}^P m_i H_K$ .

What is  $\text{OPT}(\sigma)$ : During  $i$ th and  $(i-1)$ st phases,  $\geq K + m_i$  distinct pg requests are made.

# pg faults during  $i$  &  $i-1$ 'th phase  $\geq m_i$ .

$$\begin{aligned}\therefore \text{OPT}(\sigma) &\geq \sum_{i=1}^P F_i \quad [F_i = \# \text{fault in } i\text{th phase}, \\ &\quad P = \# \text{phases}] \\ &\geq \frac{1}{2} [(F_1 + F_2) + (F_2 + F_3) + \dots] \\ &\geq \frac{1}{2} \sum_{i=1}^P m_i.\end{aligned}$$

$$\text{C.R.} \leq \sum_{i=1}^P m_i H_K / \left( \frac{1}{2} \sum_{i=1}^P m_i \right) = 2 H_K. \quad \blacksquare$$

## § Lower bounds for randomized algorithms:

Yao's lower bound principle:

- Uses L.B. for deterministic algorithms on a distribution to bound L.B. for rand algo.

$C_A(\sigma)$  be the cost of a deterministic online algorithm  $A$  with input  $\sigma$ .

Fact: Any randomized  $\mu$ -comp algo against oblivious adversaries is simply a prob. distr. over the set of deterministic algos  $\mathcal{F}$ .

$$\mathbb{E}_{A \in \mathcal{F}} \{C_A(\sigma)\} \leq \mu \cdot C_{\text{OPT}}(\sigma), \forall \sigma \quad \dots (1)$$

Consider some distr.  $D_I$  over input  $\sigma$ .

$$\mathbb{E}_{\sigma \in D_I} [\mathbb{E}_{A \in \mathcal{F}} \{C_A(\sigma)\}] \leq \mu \cdot \mathbb{E}_{\sigma \in D_I} [C_{\text{OPT}}(\sigma)]$$

Exchanging expectations (Fubini's thm):

$$\mathbb{E}_{A \in \mathcal{F}} [\mathbb{E}_{\sigma \in D_I} \{C_A(\sigma)\}] \leq \mu \cdot \mathbb{E}_{\sigma \in D_I} [C_{\text{OPT}}(\sigma)] \quad \dots (2)$$

Now,  $\min_{A_D \in \mathcal{D}} \mathbb{E}_{\sigma \in D_I} \{C_{A_D}(\sigma)\}$ , where  $\mathcal{D}$  is the set of all det. algo

$$\leq \mathbb{E}_{A \in \mathcal{F}} [\mathbb{E}_{\sigma \in D_I} \{C_A(\sigma)\}]$$

$$\leq \mu \cdot \mathbb{E}_{\sigma \in D_I} [C_{\text{OPT}}(\sigma)]$$

$$\Rightarrow \mu \geq \min_{A_D \in \mathcal{D}} \frac{\mathbb{E}_{\sigma \in D_I} [C_{A_D}(\sigma)]}{\mathbb{E}_{\sigma \in D_I} [C_{\text{OPT}}(\sigma)]} \quad \begin{array}{l} \text{Ratio of expected} \\ \text{cost of best} \\ \text{deterministic algo} \\ \& \text{the optimal} \\ \text{offline algo.} \end{array}$$

Note: Choice of  $D_I$  is arbitrary. Creative part is finding a "good"  $D_I$ .

To show L.B. for paging, we need two facts.

### Coupon Collector's Problem.

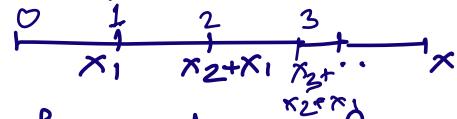
Suppose each box of cereal contains one of  $n$  different coupons. Once you obtain one of every type of coupon, you can send in for a prize.

Assuming coupon in each box is chosen independently and uniformly at random, how many boxes of cereal you need to buy before you obtain at least one of every type of coupon.

- Let  $X$  be #boxes bought until we have all types of coupons.

Let  $X_i$  denote #boxes bought while you had exactly  $(i-1)$  different coupons, then clearly

$$X = \sum_{i=1}^n X_i.$$



When exactly  $(i-1)$  coupons have been found, the prob. of obtaining a new coupon is:

$$P_i = 1 - \frac{i-1}{n}.$$

Hence,  $X_i$  is a geom RV with parameter  $P_i$ .

$$\mathbb{E}[X_i] = \frac{1}{P_i} = \frac{n}{n-i+1}.$$

\* with replacement

$$\therefore \mathbb{E}[X] = \mathbb{E}\left[\sum_{i=1}^n X_i\right] = \sum_{i=1}^n \mathbb{E}[X_i]$$

$$= \sum_{i=1}^n \frac{n}{n-i+1} = n \sum_{j=1}^n \frac{1}{j} = n H(n).$$

• **Renewal Theory :** [Section E in Borodin et al.]

A renewal process represents process that counts # times a certain event restarts itself within some time interval.

$\{X_i : i \in \mathbb{N}\}$  be sequence of +ve iid RV.

Let  $\mathbb{P}[X_i = 0] < 1$  and  $S_n = \sum_{i=1}^n X_i, n \in \mathbb{N}$

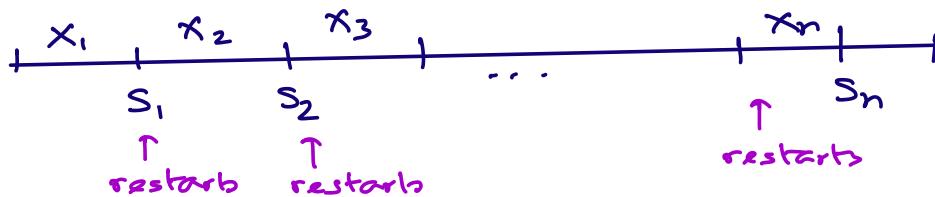
$$S_0 = 0. \quad N(t) = \sup \{n : S_n \leq t\}.$$

formally, renewal process is a stochastic process

$$\{N(t) : t \geq 0\}.$$

Remember, stochastic process  $\{x(t) : t \in T\}$  is a family of RV where  $T$  is an index set (typically time in online algo) &  $x(t)$  is the state of the process at time  $t$ .

Renewal occurs at time  $t$ , if  $S_n = t$  for some  $n$ .



After each renewal, the process begins again.

$s_n$ : time of  $n$ 'th renewal.

$$= s_n - s_{n-1}$$

$x_n$ : time between  $(n-1)$  &  $n$ 'th renewal.

$N(t)$ : Number of renewals in the interval  $[0, t]$ .

Renewal function  $M(t) = \mathbb{E}[N(t)]$ .

Elementary renewal theorem:

Given a sequence of iid RVs  $X_i, i \geq 1$ ,

with  $0 < \mathbb{E}[X_1] \leq \infty$ , then

$$\lim_{t \rightarrow \infty} \frac{M(t)}{t} = \frac{1}{\mathbb{E}[X_1]}.$$

- L.B. for randomized algo for paging is  $\ln K$ .

→ Let universe size be  $K+1$ .

Distribution  $D$ : Each input at time  $i$ ,  $\sigma_i$  is uniformly distributed over  $K+1$  elements.

Pr of fault on any online algo  $A$  on  $\sigma = \{\sigma_i\}_{i=1}^j$  for  $i \in [j]$ .

For an input  $\sigma = \{\sigma_i\}_{i=1}^j$ :

$$\min_A \mathbb{E}_D \{ F_A(\sigma) \} = \sum_{i=1}^j \frac{1}{K+1} = \frac{j}{K+1}. \quad \dots \textcircled{*}$$

# faults of  $A$  on  $\sigma$

Now we compute U.B. on  $\mathbb{E}_D \{ F_{OPT}(\sigma) \}$ .

We partition  $\sigma$  into phases  $P_i$ , where  $P_i$  consists of requests made at time  $[\theta_i, \theta_i + 1, \dots, \theta_{i+1} - 1]$ , where  $\theta_1 = 1$ , and

$$\theta_{i+1} = \min \{ r : \{\sigma_{\theta_i}, \sigma_{\theta_i+1}, \dots, \sigma_r\} = [K+1] \}.$$

Thus  $P_i$  contains exactly  $K$  distinct file requests are made.  $|P_i| = \theta_{i+1} - \theta_i$ .

$OPT = LFD$ , hence if one fault is incurred in a phase then no more faults can occur.

$$\begin{aligned} \therefore \# \text{faults incurred by } OPT \text{ until time } j \\ \leq \# \text{completed phases by time } j + 1 \end{aligned}$$

$$\text{Hence, } \mathbb{E}[F_{OPT}(\sigma)] \leq 1 + \mathbb{E}[\max \{ i : \theta_i \leq j \}]. \quad \textcircled{1}$$

Now  $D$  is uniform & iid,  $|P_i| = \theta_{i+1} - \theta_i$  is also iid.

So apply elementary renewal theorem,  
with renewal intervals.  $X_i = P_i = \theta_{i+1} - \theta_i$ ,

$$\lim_{j \rightarrow \infty} (1 + \mathbb{E}[\max\{e : \theta_e \leq j\}]) / j = 1 / \mathbb{E}[P_1]. \quad (2)$$

$\approx$  When  $|S| \rightarrow \infty$ , the number of completed phases normalized with  $|S|$  converges to the reciprocal of the expected length of a phase.

$$\lim_{j \rightarrow \infty} \frac{F_{\text{OPT}}(j)}{j} \leq 1 / \mathbb{E}[P_1]. \quad \dots \quad (3)$$

From defn of phase,  $P_1$  is identically distr. as the finish time  $T_c$  in the coupon collector problem.

$$\begin{aligned} \text{Hence, } \mathbb{E}(P_1) &= (K+1) H_K - 1 \\ &= (K+1) H_K + 1 - 1 \\ &= (K+1) H_K. \end{aligned} \quad \dots \quad (4)$$

$$\text{Therefore, } \mathbb{E}[F_{\text{OPT}}(j)] \leq j / \mathbb{E}[P_1] \leq \frac{j}{(K+1) H_K}$$

$$\text{Also, } \min_{\sigma \in D} \mathbb{E}_{\sigma} \{ F_{\sigma}(j) \} = \frac{j}{K+1}. \quad \cdot \quad (\dagger\dagger)$$

$$\begin{aligned} \therefore \mu &\geq \min_{A \in D} \frac{\mathbb{E}_{\sigma \in D} [C_{A\sigma}(j)]}{\mathbb{E}_{\sigma \in D} [C_{\text{OPT}}(\sigma)]} \\ &\stackrel{(\text{from } \dagger \& \dagger)}{\geq} \frac{j/K+1}{j/(K+1) H_K} \geq H_K \geq \ln K. \quad \blacksquare \end{aligned}$$